


Intermediate Level Introduction to Computing at CARC

With Machine Learning Examples

Matthew Fricke

Version 0.1



Goals

- 1) SLURM scheduler literacy
- 2) Run Machine Learning Tools on CARC Systems
 - Random Forests
 - Support Vector Machine
 - Multi-Layer Perceptron
 - Convolutional Neural Network
- We wont cover file transfer, storage systems, module system, conda, PBS. (These are all covered in depth in the video tutorials)

Logging into Hopper



First login to the Linux **workstation** in front of you. Your CARC username is on the sign in sheet.

If you have logged in before use your **existing password**

Otherwise, your initial password is **Welcome2CARC**

This is an “important step” so don’t let me move on until you have logged in

Logging into Hopper



```
ssh vanilla@hopper.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Replace vanilla with your name (unless your last name is Ice)

Logging into Hopper

Welcome to Hopper

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to <http://carc.unm.edu>, select the "New Users" menu and then click "Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the `chmod` command to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing, supported in part by the National Science Foundation, for providing the research computing resources used in this work."

Please send citations to publications@carc.unm.edu.

There are three types of slurm partitions on Hopper:

1) General - this partition is accessible by all CARC users.

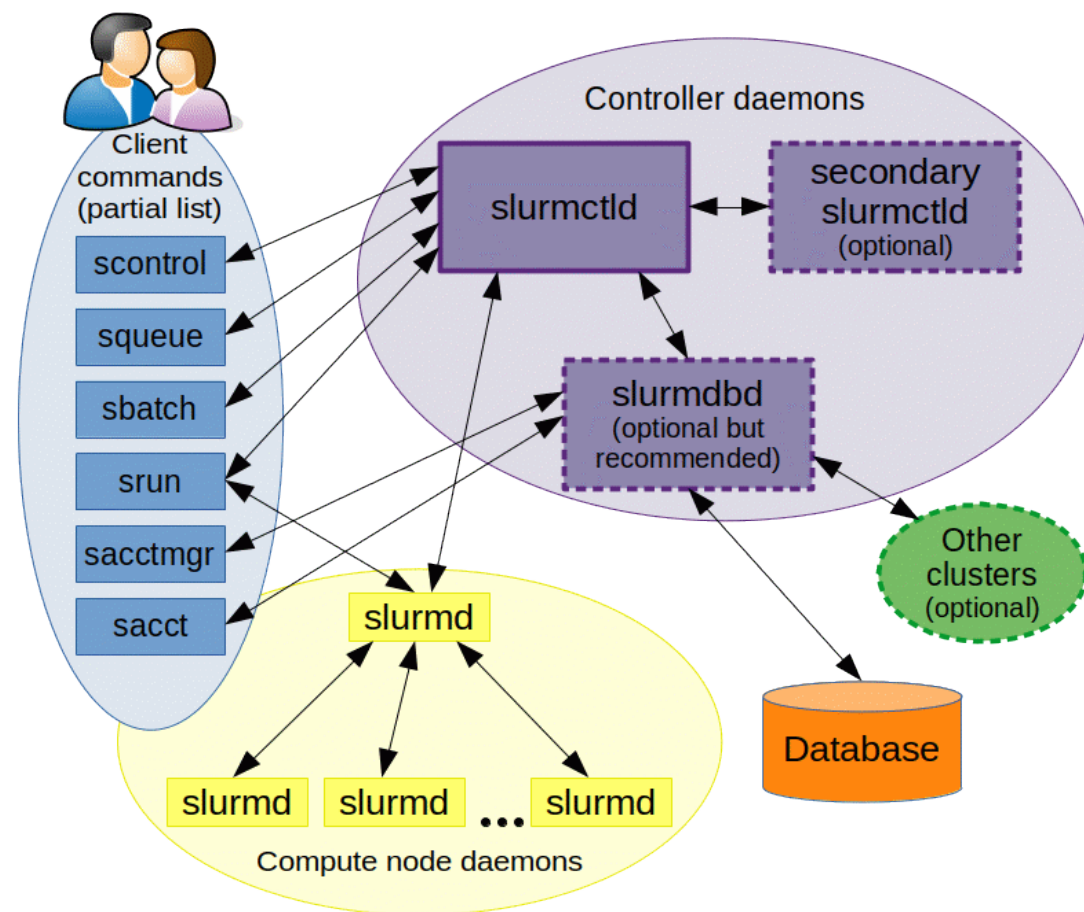
2) Condo - preemptable scavenger queue available to all condo users. Your job must use checkpointing to use this queue or you will lose any work you have done if it is preempted by the partition's owner.

3) Named partitions - these partitions are available to condo users working under the grant/lab/center that purchased the associated hardware.


Type "qgrok" to get the status of the partitions.

Last login: Wed Jul 27 17:46:13 2022 from 129.24.246.68
mfricke@hopper:~ \$

Simple Linux Utility for Resource Management





VOTED  #1 SOFT DRINK OF THE 31ST CENTURY!



[vanilla@hopper ~]\$ qgrok

queues	free	busy	offline	jobs	nodes	CPU	GPU	CPU/node	GPU/node	Memory/node	time_limit	CPU_limit
general	4	6	0	4	10	320	0	32	0	93G	2-00:00:00	64
debug	2	0	0	0	2	64	0	32	0	93G	4:00:00	8
condo	22	25	4	7	51	1632	28	32	2	93G-1.5T	2-00:00:00	192
bugs	2	0	0	0	2	64	0	32	0	93G	7-00:00:00	

vanilla@hopper:~ \$ qgrok

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
--------	------	------	---------	------	-------	------	------

general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1	0	0	2	64	0
pathogen	1	0	0	0	1	32	0
tc	5	5	0	3	10	32	0
gold	2	0	0	0	2	64	0
fishgen	0	1	0	0	1	32	0
neuro-hsc	8	6	0	0	14	448	0
cup-ecs	0	2	0	2	2	64	0
tid	0	1	0	0	1	32	0
biocomp	0	1	0	0	1	32	0
chakra	1	0	0	0	1	32	0
pna	0	0	1	0	1	32	0
totals:	19	28	1	14	48	1536	8

Open partitions for use by everyone with a CARC account.

Purchased by the Office for the Vice President for Research.

vanilla@hopper:~ \$ qgrok

queues free busy offline jobs nodes CPUs GPUs

-----	-----	-----	-----	-----	-----	-----	-----
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1	0	0	2	64	0
pathogen	1	0	0	0	1	32	0
tc	5	5	0	3	10	320	0
gold	2	0	0	0	2	64	0
fishgen	0	1	0	0	1	32	0
neuro-hsc	8	6	0	0	14	448	0
cup-ecs	0	2	0	2	2	64	4
tid	0	1	0	0	1	32	2
biocomp	0	1	0	0	1	32	1
chakra	1	0	0	0	1	32	1
pna	0	0	1	0	1	32	0
totals:	19	28	1	14	48	1536	8

vanilla@hopper:~ \$ qgrok

queues free busy offline jobs nodes CPUs GPUs

```
-----
general  1   9   0   7  10 320 0
debug    2   0   0   0   2  64 0
condo    18  19   1   7  38 1216 8
bugs     0   2   0   0   2  64 0
pcnc     1   1   0   0   2  64 0
pathogen 1   0   0   0   1  32 0
tc       5   5   0   3  10 32 0
gold     2   0   0   0   2  64 0
fishgen  0   1   0   0   1  32 0
neuro-hsc 8   6   0   0  14 32 0
cup-ecs  0   2   0   2   2  64 0
tid      0   1   0   0   1  32 0
biocomp  0   1   0   0   1  32 0
chakra   1   0   0   0   1  32 0
pna      0   0   1   0   1  32 0
totals:  19  28   1  14  48 1216 8
```

Private partitions

- Reserved for use by the purchaser.
- Request access by emailing support@carc.unm.edu and CC the partition owner.

```
mfricke@hopper:~$ qgrok
```

```
queues  free  busy  offline  jobs  nodes  CPUs  GPUs
```

```
-----
general  1   9   0   7  10  320  0
debug    2   0   0   0   2   64  0
condo    18  19   1   7  38 1216  8
bugs     0   2   0   0   2   64  0
pcnc     1   1   0   0   2   64  0
pathogen 1   0   0   0   1   32  0
tc       5   5   0   3  10  32  0
gold     2   0   0   0   2   64  0
fishgen  0   1   0   0   1   32  0
neuro-hsc 8   6   0   0  14  32  0
cup-ecs  0   2   0   2   2   64  0
tid      0   1   0   0   1   32  0
biocomp  0   1   0   0   1   32  0
chakra   1   0   0   0   1   32  0
pna      0   0   1   0   1   32  0
totals:  19  28   1  14  48 1216  8
```

Condo “scavenger” partition

- Allows you to use compute nodes purchased by another group that are currently idle.
- May be interrupted at any time if the owners start to use it.

[vanilla@hopper ~]\$ **quotas**

Home Directory (/users/vanilla):

quota: Cannot resolve mountpoint path /root/.spack: Permission denied

Disk quotas for user vanilla (uid 659):

Filesystem	space	quota	limit	grace	files	quota	limit	grace
chama:/home/homes	1527M	100G	200G		14913	4295m	4295m	

Centerwide user scratch (/carc/scratch/users/mfricke)

Quota information for storage pool Default (ID: 1):

user/group		size		chunk files	
name	id	used	hard	used	hard
-----	-----	-----	-----	-----	-----
mfricke	1512	592.71 GiB	1024.00 GiB	32784	unlimited

Centerwide scratch quota for project mfricke2016174 (/carc/scratch/projects/mfricke2016174)

Quota information for storage pool Default (ID: 1):

user/group		size		chunk files	
name	id	used	hard	used	hard
-----	-----	-----	-----	-----	-----
mfricke2016174	2016142	190.97 GiB	1024.00 GiB	23704	unlimited

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up    4:00:00    2  idle  hopper[011-012]
```

sinfo reports information about
partitions


```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST  
debug      up    4:00:00    2  idle hopper[011-012]
```

The debug queues are intended
for testing your programs.

And for interactive jobs.

```
[vanilla@hopper ~]$ sinfo --partition debug
```

```
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
```

```
debug      up    4:00:00    2 idle hopper[011-012]
```



Name

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up      4:00:00    2  idle hopper[011-012]
```



You can run a “job” for up to 4 hrs.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST  
debug      up    4:00:00    2  idle hopper[011-012]
```



There are two nodes in this partition.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST
debug      up    4:00:00    2  idle hopper[011-012]
```



The names of the nodes in the
partition

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST
debug      up    4:00:00    2  idle hopper[011-012]
```



The names of the nodes in the
partition


```
[vanilla@hopper ~]$ sinfo --partition general
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
general*	up	2-00:00:00	9	alloc	hopper[001-009]
general*	up	2-00:00:00	1	idle	hopper010



Hopper001 through 009 are running jobs.

Hopper010 is waiting to be used.

```
[vanilla@hopper ~]$ hostname  
hopper  
[vanilla@hopper ~]$
```



Running on the Head Node.

The head node's name is "hopper".

```
[vanilla@hopper ~]$ hostname
```

```
hopper
```

```
[vanilla@hopper ~]$ man hostname
```

```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man man  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man sinfo
```

sinfo(1)

Slurm Commands

sinfo(1)

NAME

sinfo - View information about Slurm nodes and partitions.

SYNOPSIS

sinfo [OPTIONS...]

DESCRIPTION

sinfo is used to view partition and node information for a system running Slurm

OPTIONS

-a, --all

Display information about all partitions. This causes information to be displayed about partitions that are configured as hidden and partitions that are unavailable to the user's group.

[vanilla@hopper ~]\$ sinfo --all

```
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST
general*   up 2-00:00:00    9 alloc hopper[001-009]
general*   up 2-00:00:00    1 idle hopper010
debug      up  4:00:00     2 idle hopper[011-012]
condo      up 2-00:00:00     1 down* hopper045
condo      up 2-00:00:00     3  mix hopper[018-020]
condo      up 2-00:00:00    16 alloc hopper[013-015,028-036,049-052]
condo      up 2-00:00:00    18 idle hopper[016-017,021-027,037-044,053]
bugs       up 7-00:00:00     2 alloc hopper[013-014]
pcnc       up 7-00:00:00     1 alloc hopper015
pcnc       up 7-00:00:00     1 idle hopper016
pathogen   up 7-00:00:00     1 idle hopper017
tc         up 7-00:00:00     3  mix hopper[018-020]
tc         up 7-00:00:00     2 alloc hopper[029-030]
tc         up 7-00:00:00     5 idle hopper[021-025]
gold       up 7-00:00:00     2 idle hopper[026-027]
fishgen    up 7-00:00:00     1 alloc hopper028
neuro-hsc  up 7-00:00:00     6 alloc hopper[031-036]
neuro-hsc  up 7-00:00:00     8 idle hopper[037-044]
cup-ecs    up 7-00:00:00     2 alloc hopper[049-050]
tid        up 7-00:00:00     1 alloc hopper051
biocomp    up 7-00:00:00     1 alloc hopper052
chakra     up 7-00:00:00     1 idle hopper053
pna        up 7-00:00:00     1 down* hopper045
```



```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Tell slurm to run a program
on a compute node...

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Run the program on a
compute node in the
debug partition.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



The program
to run.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
[vanilla@hopper ~]$ queue
```

```
[vanilla@hopper ~]$ queue
```

**PD means programs
that are waiting their
turn.**

that are waiting their turn.

Shows you what the slurm scheduler is doing right now.

Here we can see that user 'erowland' has a lot of programs waiting to run.

JOBID	PARTITION	NAME	USER	ST	TIME
4314	general	PRE	erowland	PD	0:00
4315	general	PRE	erowland	PD	0:00
4317	general	PRE	erowland	PD	0:00
4318	general	PRE	erowland	PD	0:00

Shows you what the slurm scheduler is doing right now.

Here we can see that user 'erowland' has a lot of programs waiting to run.


```
[vanilla@hopper ~]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(Reason)	
4314	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4315	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4317	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4318	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4319	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4320	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4321	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4322	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4323	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4324	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4325	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4326	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4327	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4328	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4329	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4330	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4331	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4332	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4333	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4334	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4335	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4336	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	
4337	general		erowland	PD	0:00	2	(QOSMaxCpuPerUserLimit)	

The reason these jobs are not running is that 'erowland' is already using the maximum number of CPUs they are allowed.

```
[vanilla@hopper ~]$ squeue -t R --all
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
4405	condo	2ndMA	mfricke	R	1-07:48:30	6	hopper[031-036]
5208	condo	NN	kgu	R	5:48:49	1	hopper015
5210	condo	NN	kgu	R	6:30:13	1	hopper014
5209	condo	NN	kgu	R	6:31:13	1	hopper013
5206	condo	NN	kgu	R	6:32:13	1	hopper051
5207	condo	NN	kgu	R	6:32:13	1	hopper052
5205	condo	NN	kgu	R	6:32:43	1	hopper028
4595	cup-ecs	golConfi	aalasand	R	2-06:51:59	1	hopper050
4594	cup-ecs	golConfi	aalasand	R	2-06:52:03	1	hopper049
5120	general	jupyterh	jacobm	R	11:45:47	1	hopper007
4313	general	PRE	erowland	R	1:17:29	2	hopper[003-004]
5111	general	1stMA	mfricke	R	11:15:28	2	hopper[005-006]
5025	general	c2n	jxzuo	R	1:50	1	hopper001
5024	general	c2n	jxzuo	R	31:28	1	hopper002
5203	general	NN	kgu	R	6:37:50	1	hopper009
5201	general	NN	kgu	R	6:38:14	1	hopper008
4390	tc	UCsTpCyd	lepluart	R	2-15:18:18	3	hopper[018-020]
5198	tc	NN	kgu	R	6:40:19	1	hopper030
5196	tc	NN	kgu	R	6:40:31	1	hopper029

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

You ran two **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

You ran eight **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

By default, each task (copy of your program) is allowed to use one CPU.

Many programs are able to use more than one CPU at a time.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.  
hopper011  
hopper011
```

**Here we are telling SLURM to
run 2 copies of our program
and let each copy of our
program use 2 CPUs.**

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2 --ntasks-per-node 4 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
hopper012
```

You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.

```
hopper012  
hopper011  
hopper011  
hopper012  
hopper012  
hopper011  
hopper011
```

Here we are telling SLURM to run 4 copies of our program on 2 different compute nodes.

This is useful when our programs need a bigger share of the compute node.


```
[vanilla@hopper ~]$ srun --partition debug --nodes 2  
--ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your  
submission script.
```

```
hopper011
```

```
hopper012
```

```
hopper012
```

And we can combine all three.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been allocated  
submission script.
```

```
hopper011
```

```
Hopper011
```

**And we can specify how much
memory we want.**

**--mem 4G means give me 4
gigabytes of memory per node.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been allocated  
submission script.
```

```
hopper011
```

```
Hopper011
```

Why does all this matter?

**The purpose of SLURM is to provide
you the hardware your programs
need.**

**So you have to understand what
those requirements are really well.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been all  
submission script.
```

```
hopper011
```

```
Hopper011
```

- 1) Can my program use multiple CPUs?
- 2) How much memory does my program need?
- 3) Can my program use multiple compute nodes (MPI*, GNU Parallel*)?
- 4) Can my program use GPUs?

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project  
hopper012  
hopper012  
You have not been allocated a  
submission script.  
hopper011  
Hopper011
```

This command is getting pretty long.

We can use **shell scripts** to automate
all this in **batch mode**.

Interactive vs Batch Mode

Interactive Mode

- Everything so far has been interactive. You request hardware, run your program, and get the output on your screen right away.

Batch Mode

- Most programs at an HPC center are run in “batch” mode.
- Batch mode means we write a shell script that the SLURM scheduler runs for us. The script requests hardware just like we did with `salloc` and then runs the commands in the script.
- Whatever would have been written to the screen is saved to a file instead.

```
[vanilla@hopper ~]$ git clone https://lobogit.unm.edu/CARC/workshops.git
```

```
Cloning into 'workshops'...
```

```
remote: Enumerating objects: 132, done.
```

```
remote: Counting objects: 100% (75/75), done.
```

```
remote: Compressing objects: 100% (43/43), done.
```

```
remote: Total 132 (delta 33), reused 74 (delta 32), pack-reused 57
```

```
Receiving objects: 100% (132/132), 57.58 KiB | 3.60 MiB/s, done.
```

```
Resolving deltas: 100% (51/51), done.
```

Rather than make you write shell scripts lets just download some we wrote for this workshop...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H2O.gjf
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run tree to see how the workshops directories are organized...


```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H2O.gif
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run **tree** to see how the workshops directories are organized...

The workshop files are divided into “code”, “slurm”, and “data” directories.

```
[vanilla@hopper intro_workshop]$ pwd  
/users/vanilla/workshops/intro_workshop  
[vanilla@hopper intro_workshop]$ cat slurm/workshop_example1.sh  
#!/bin/bash  
#SBATCH --partition debug  
#SBATCH --ntasks 4  
#SBATCH --time 00:05:00  
#SBATCH --job-name ws_example  
#SBATCH --mail-user your_username@unm.edu  
#SBATCH --mail-type ALL  
  
srun hostname
```

Let's take a look at the **workshop_example.sh** script in the slurm directory...

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh  
sbatch: Account not specified in script or ~/.default_slurm_account, using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

We **submit** our slurm
shell script with the
sbatch command.

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh
sbatch: Account not specified in script or ~/.default_slurm_account, using latest project
Submitted batch job 5252
[vanilla@hopper intro_workshop]$
```

Notice that the only output we get is a job id.

This indicates that the script was successfully sent to the scheduler.

The commands in the script will run as soon as the hardware requested is available.

We **submit** our slurm shell script with the sbatch command.

Workflow

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Compute Node 01

Compute Node 02

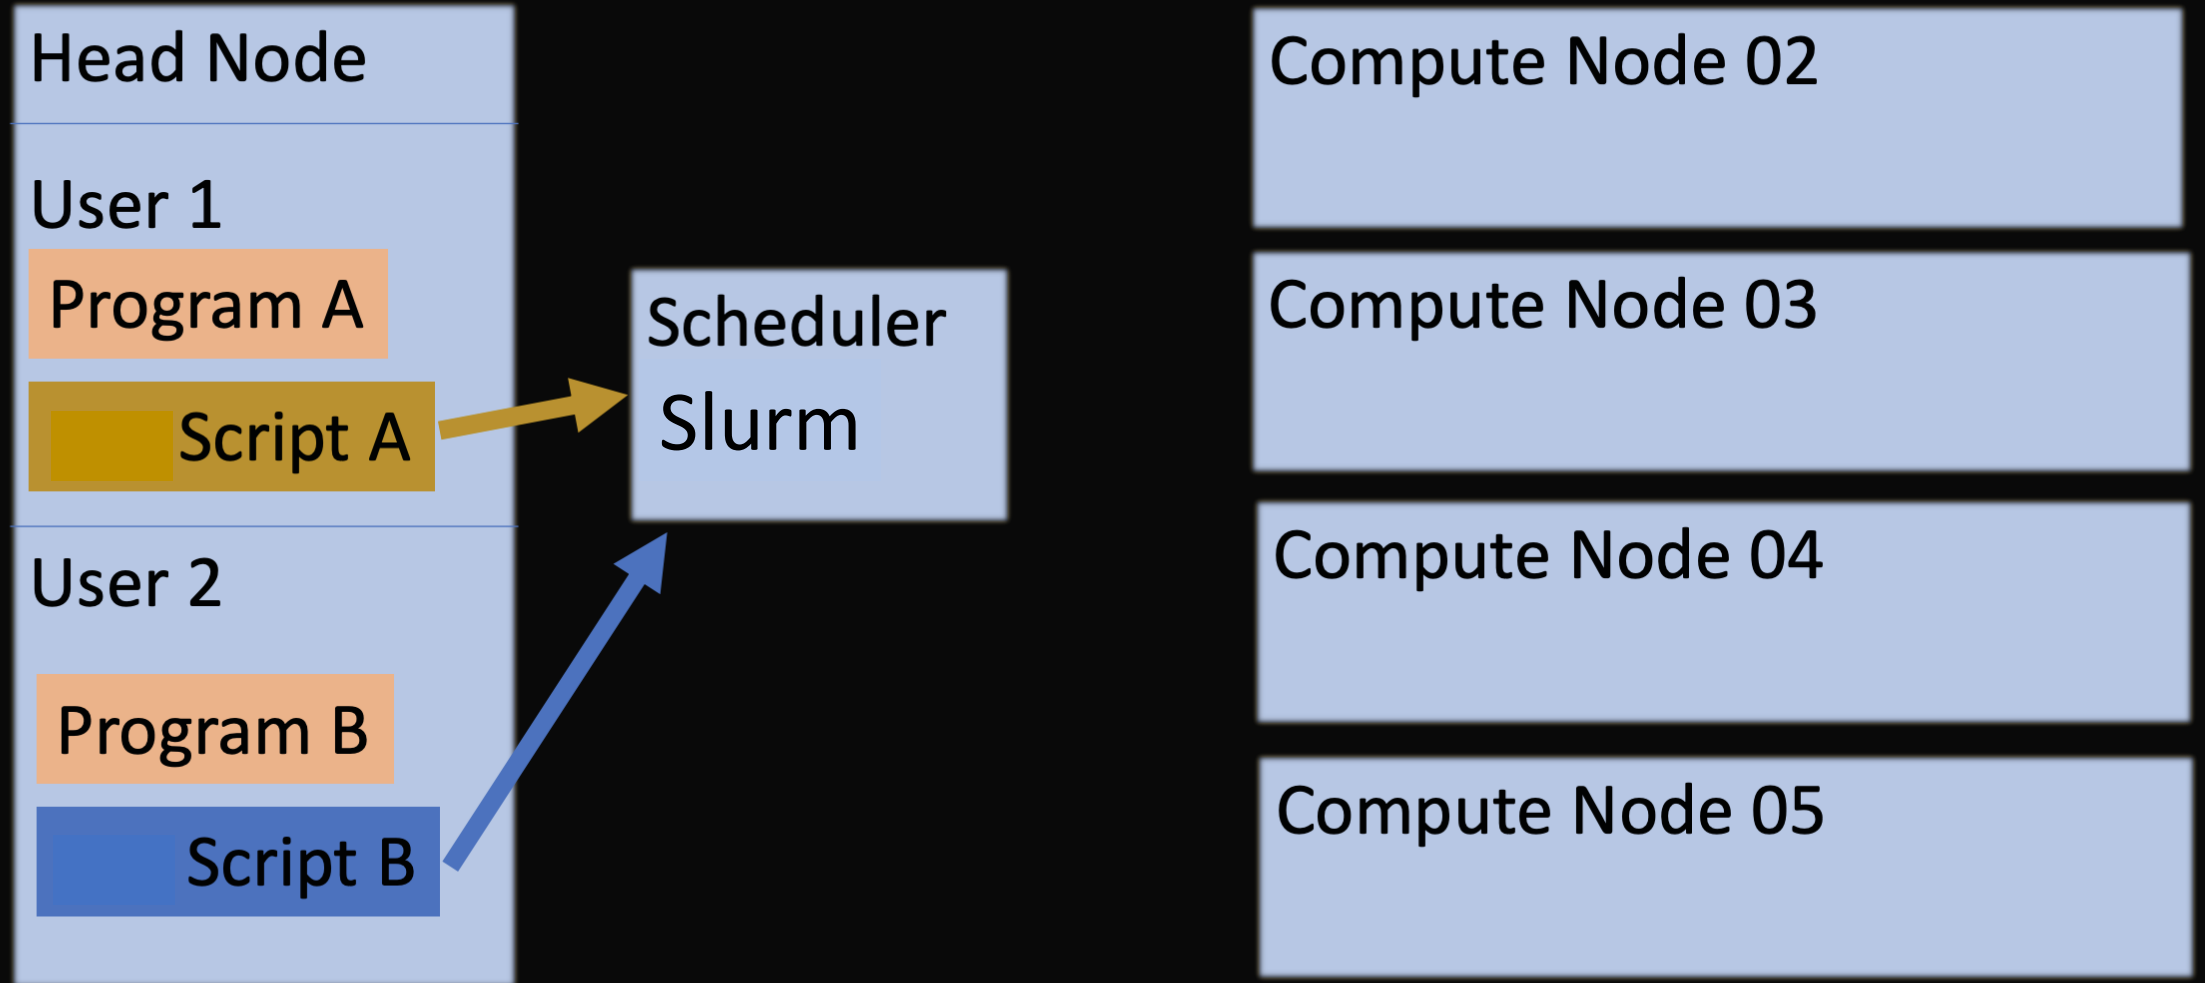
Compute Node 03

Compute Node 04

Compute Node 05

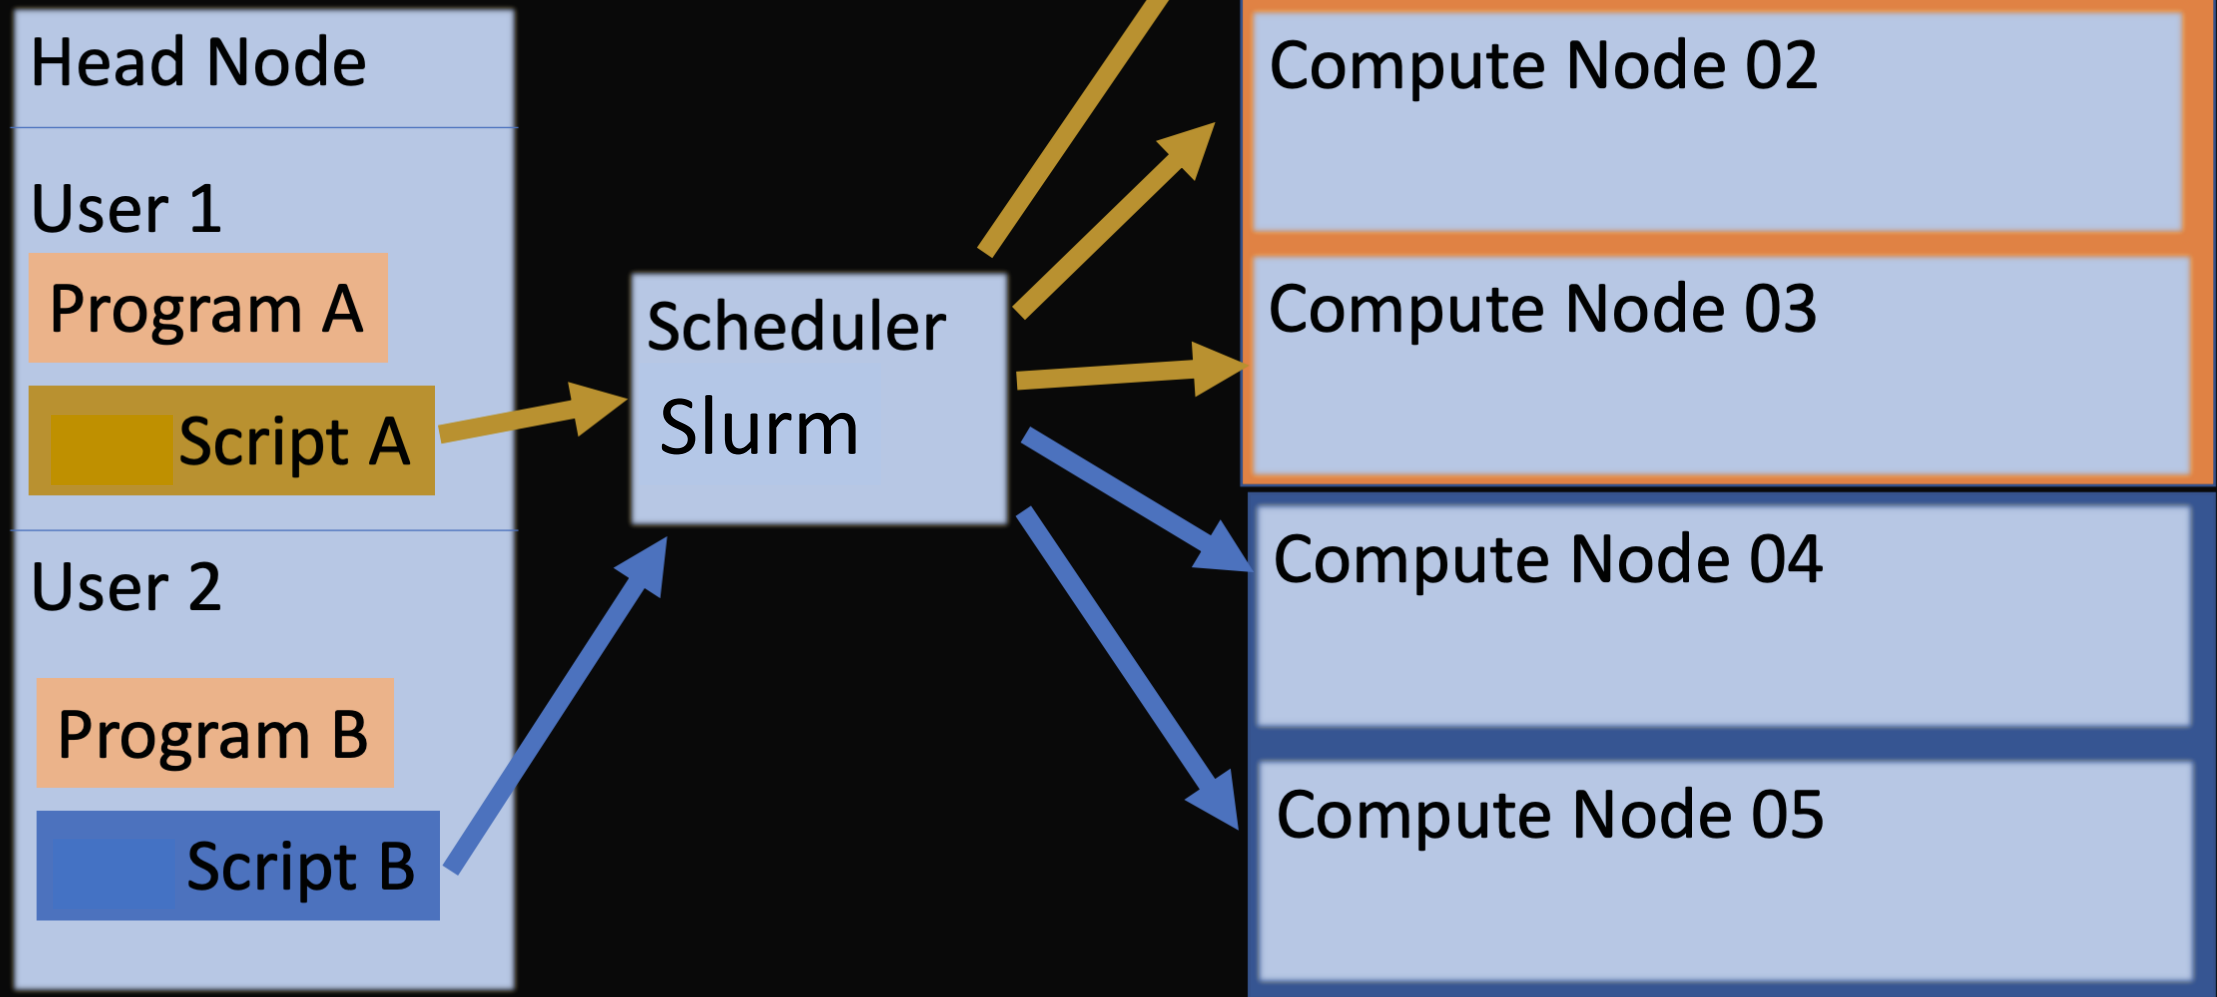
Shared filesystems – All nodes can access the same programs and write output

Workflow



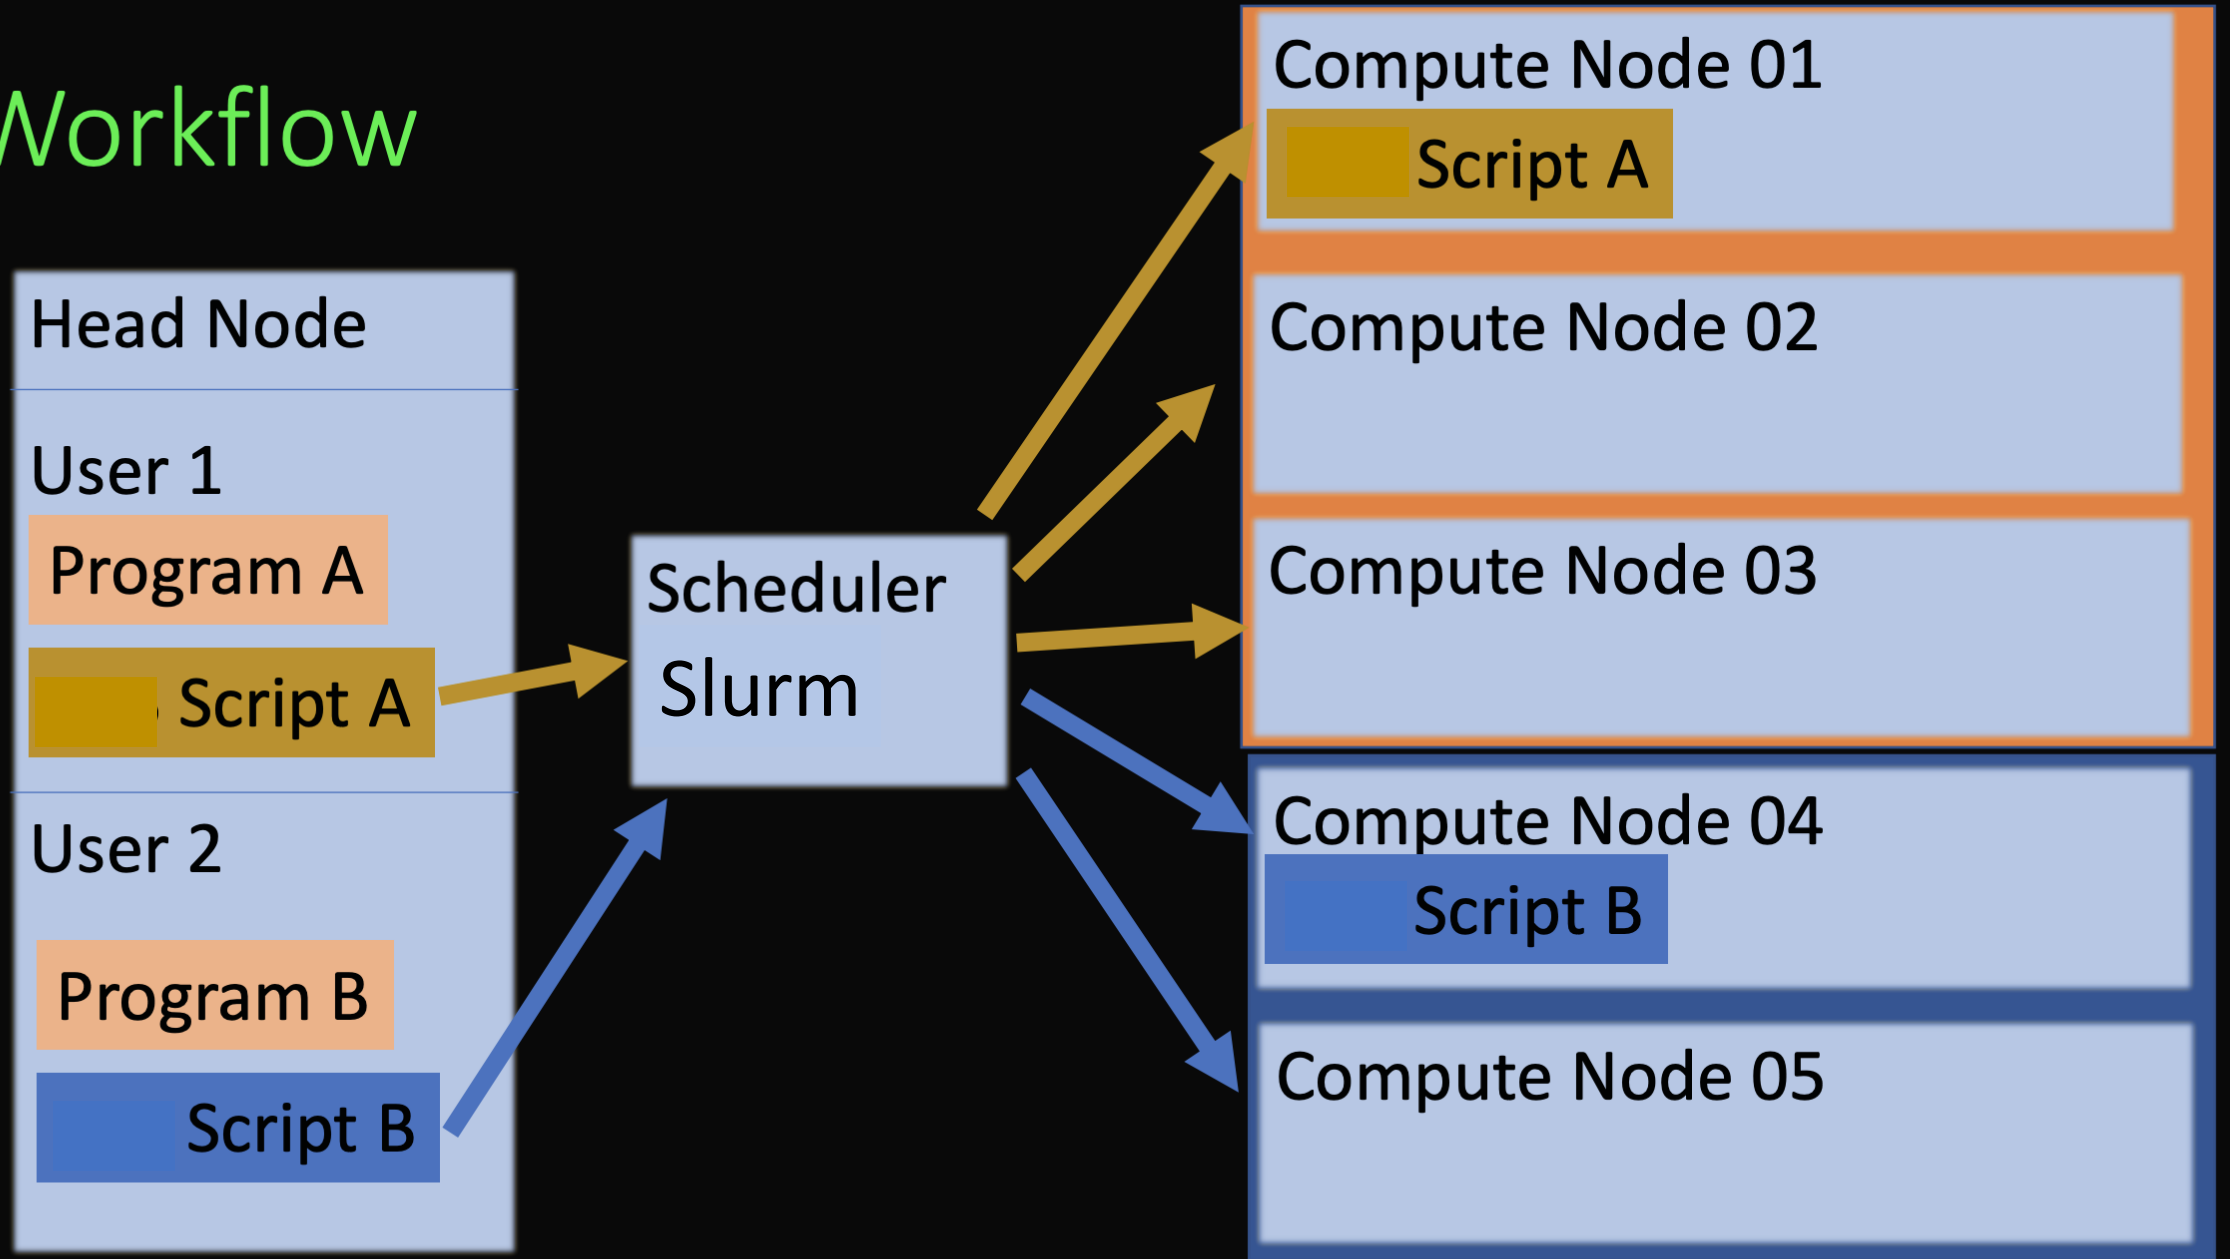
Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

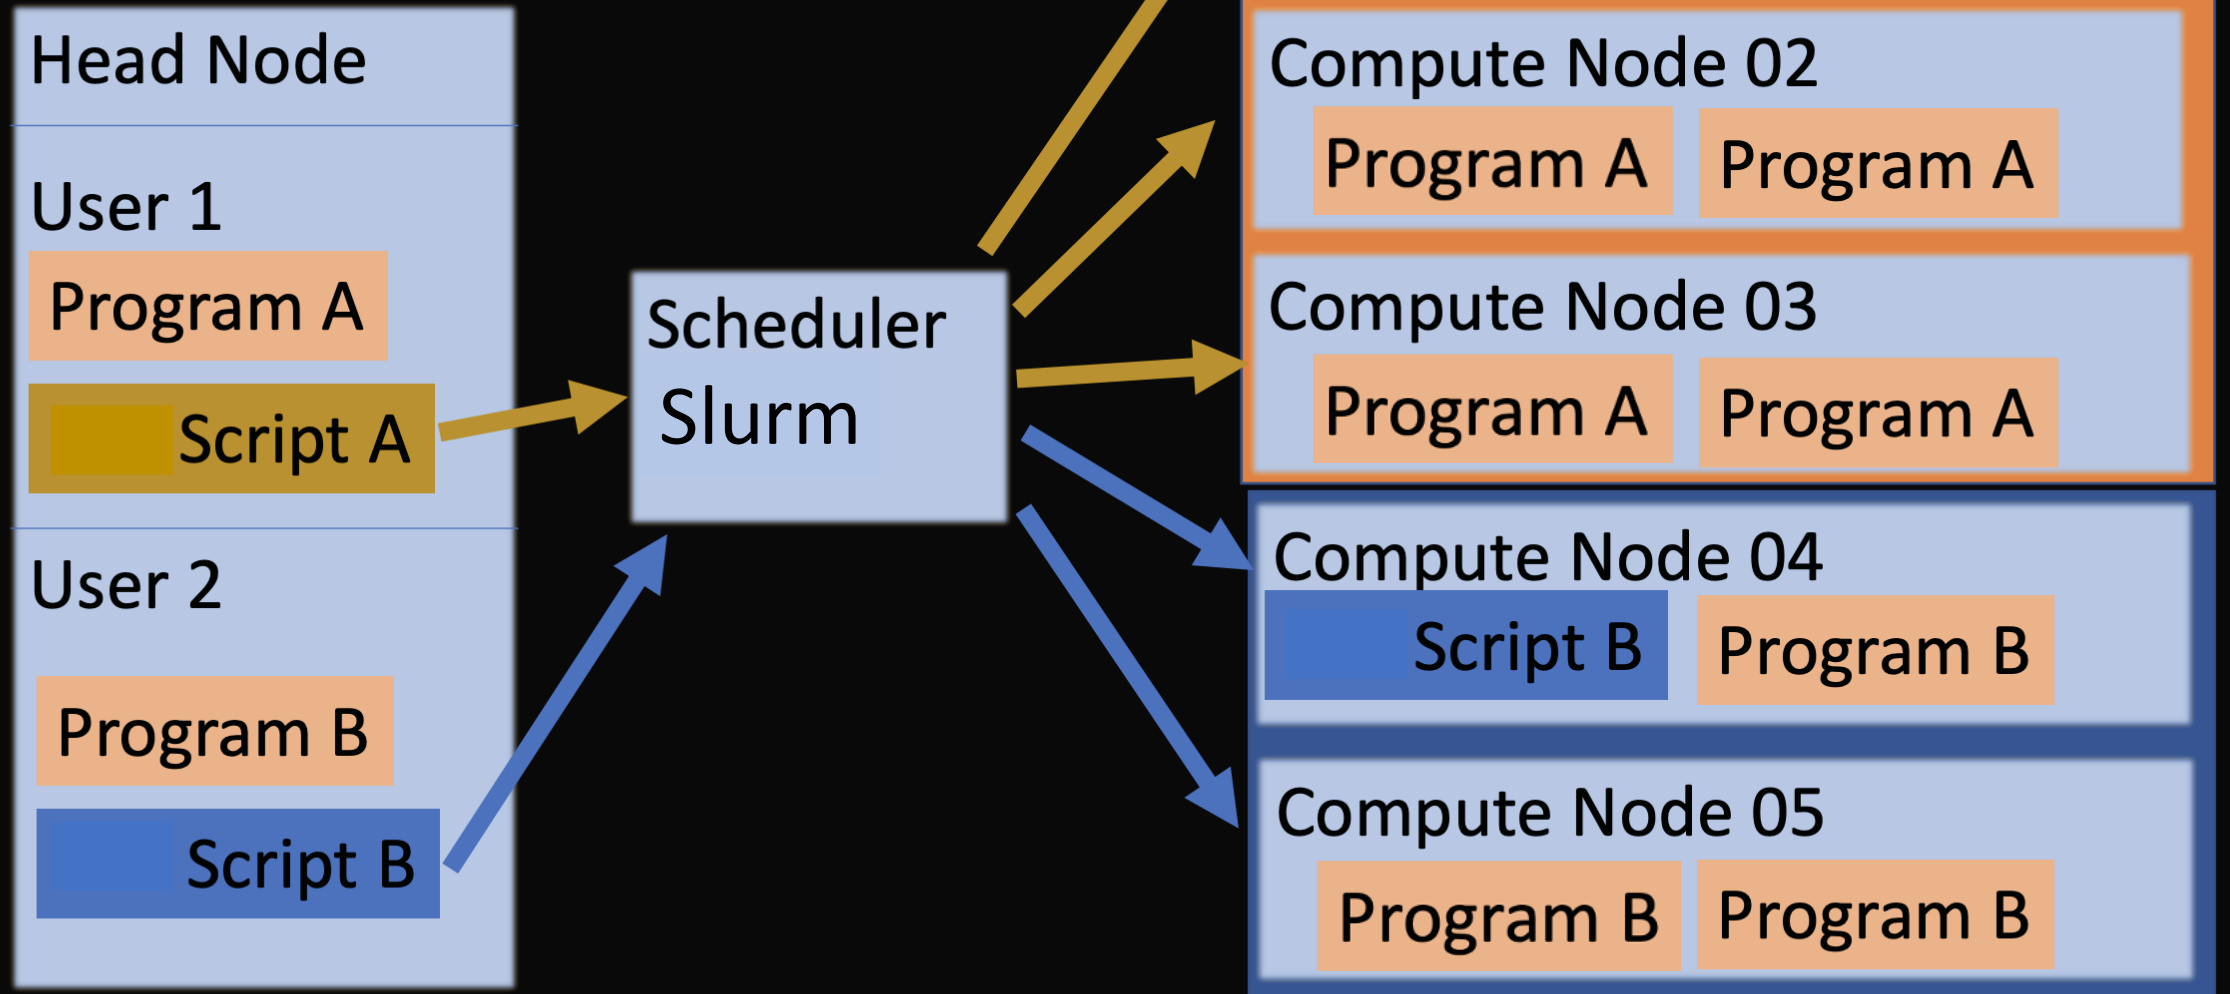
Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow


We need something in the script to run the program on all the nodes. E.g. srun.



Shared filesystems – All nodes can access the same programs and write output

```
[vanilla@hopper intro_workshop]$ ls  
code data pbs slurm slurm-5252.out
```

The **hostname**
command is very fast
so everyone's job
should finish in a few
seconds.



When it is finished you will
have a new file named
slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ ls  
code data pbs slurm slurm-5252.out
```



When it is finished you will have a new file named `slurm-{your job id}.out`.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011
```

```
[vanilla@hopper intro_workshop]$ module load miniconda3
```

```
[vanilla@hopper intro_workshop]$ conda create -n numpy numpy
```

Wait a while – introduce yourselves to your neighbor...

Conda allows you to install software into your home directory. In this case we need the numerical python libraries for calcPiSerial.py

Let's experiment with a program that does slightly more than print the hostname.

```
[vanilla@hopper intro_workshop]$source activate numpy
```

```
[vanilla@hopper intro_workshop]$srun --partition debug python code/calcPiSerial.py 10
```

```
srun: Using account 2016199 from ~/.default_slurm_account
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
Pi = 3.14242598500109870940, (Diff=0.00083333141130559341) (calculated in 0.000005 secs with 10 steps)
```

**Activate the numpy environment and
Run calcPiSerial.py on a compute node.**

**For our example program the more steps it takes the
more accurate it is, but the longer it takes.**

```
[vanilla@hopper intro_workshop]$ cat slurm/calc_pi_serial.sh
#!/bin/bash
#SBATCH --partition debug
#SBATCH --ntasks 1
#SBATCH --time 00:05:00
#SBATCH --job-name calc_pi_serial
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

module load miniconda3
source activate numpy

cd $SLURM_SUBMIT_DIR
python code/calcPiSerial.py 1000000000
[vanilla@hopper intro_workshop]$
```

```
[vanilla@hopper intro_workshop]$ source deactivate  
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh  
sbatch: Using account 2016199 from ~/.default_slurm_account
```

Submitted batch job 5263

```
vanilla@hopper:~/workshops/intro_workshop$ squeue --me
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
5263	debug	calc_pi_	vanilla	R	0:44	1	hopper011

Edit `slurm/calc_pi_serial.sh`.

Change the email address to your address and submit the script.

Then enter `squeue --me` to see the job status.

Take a look at the job output.

Install the tools we need

- Tensorflow – Machine Learning from Google
- Scikit-learn – Open Source (was a Google summer intern project)
- Ipykernel – So we can use our python libraries in a Jupyter Notebook
- Matplotlib and Seaborn – For plotting and graphics
- Pandas – Data Science
- Numpy – Numerical libraries

And lots of other things will be installed...


```
[vanilla@hopper intro_workshop]$ conda create --name  
machine_learning matplotlib ipykernel scikit-learn  
seaborn tensorflow-gpu=2.4.1 tqdm --channel defaults
```

<snip>

```
zipp          pkgs/main/linux-64::zipp-3.17.0-py39h06a4308_0  
zlib          pkgs/main/linux-64::zlib-1.2.13-h5eee18b_1  
zstd          pkgs/main/linux-64::zstd-1.5.5-hc292b87_2
```

Proceed ([y]/n)?



Don't forget to enter "y" here

Using conda environments

- Conda create
- Conda delete
- `conda info --envs`

Conda activate {env name}

```
conda create --name machine_learning matplotlib  
ipykernel scikit-learn seaborn tensorflow-  
gpu=2.4.1 tqdm --channel defaults
```

Sign in

Username:

Password:

Sign in

Server Options

Select a job profile:

Debug Queue, 1 hours, 1 core, 4GB RAM



Start

You might not see this the first time.

hopper.alliance.unm.edu

jupyterhub

Logout

Control Panel

Files

Running

Clusters

Select items to perform actions on them.

Upload

New

☐

0

/ workshops / machine_learning / jupyter_notebooks

Name

Last Modified

File size

..

☐

classifier_examples.ipynb

Running

seconds ago

39.3 kB

Jupyter Notebook



The conda environment we installed

CARC Workshop Machine Learning

Random Forest Classifier of Penguins

For the random forest example we will read in a csv text file that contains information about penguins. The label is the species and the data to map to species are location, beak and flipper dimensions and weight. There are three species of penguin in the dataset: Adelie, Gentoo, and Chinstrap.

```
In [ ]: 1 # We will be using matplotlib for graphics throughout
        2 import matplotlib.pyplot as plt
        3
```



Logout

Control Panel

File Edit View Insert Cell Kernel Widgets Help

Trusted

Python [conda env:.conda-machine_learning_xena]

New Notebook

Open...

Make a Copy...

Save as...

Rename...

Save and Checkpoint

Revert to Checkpoint

Print Preview

Download as

Trusted Notebook

Close and Halt

AsciiDoc (.asciidoc)

HTML (.html)

LaTeX (.tex)

Markdown (.md)

Notebook (.ipynb)

PDF via LaTeX (.pdf)

reST (.rst)

Python (.py)

Reveal.js slides (.slides.html)

PDF via pyppteer (.html)

Download the notebook as a python file



C Workshop Machine Learning

om Forest Classifier of Penguins

andom forest example we will read in a csv text file that contains information about penguins. The label is the species and the data to map to re location, beak and flipper dimensions and weight. There are three species of penguin in the dataset: Adelie, Gentoo, and Chinstrap.

ib for graphics throughout
plt

ibes penguin species into a pandas dataframe
projects/share

/penguins.csv")

In []:

```
1 # C
2 # s
3 lab
4
5 classes = penguin_df[label].unique().tolist()
6 print(f"Label classes: {classes}")
7
8 penguin_df[label] = penguin_df[label].map(classes.index)
```

```
[vanilla@hopper machine_learning]$ sbatch slurm/ml_example_hopper.sh
sbatch: Account not specified in script or ~/.default_slurm_account, using latest project
Submitted batch job 5252
```

```
[vanilla@hopper machine_learning]$ tail -f slurm-{job ID}.out
```

```
32 by 32 by 3
[1]
Train data size: 50000
Train target size: 50000
Train data size: 26032
Train data size: 26032
100%|██████████| 50000/50000 [01:09<00:00, 723.6]
25%|███| 6431/26032 [00:08<00:26, 726.59]
```

**We can monitor the
output live**


```
[vanilla@xena machine_learning]$ tail -f slurm-687248.out
```

```
<snip>
```

```
665/665 [=====] - 3s 4ms/step - loss: 0.2996 - accuracy: 0.9135 - val_loss: 0.4952 - Epoch 00014:  
val_accuracy improved from 0.86947 to 0.87267, saving model to checkpoints_best_cnn/checkpoint.weights.h5
```

```
Epoch 15/30
```

```
665/665 [=====] - 3s 4ms/step - loss: 0.2996 - accuracy: 0.9135 - val_loss: 0.4952 - val_accuracy:  
0.8727
```

```
[1587 5711 3958 2686 2723 2270 2499 1981 1039 1578]
```

```
[[1528 30 32 22 6 9 46 24 12 35]
```

```
[ 68 4617 86 75 81 19 14 105 26 8]
```

```
[ 21 38 3815 71 51 19 8 67 33 26]
```

```
[ 30 78 113 2236 23 129 25 40 74 134]
```

```
[ 27 93 47 52 2183 22 24 20 27 28]
```

```
[ 10 13 31 110 21 2021 106 13 20 39]
```

```
[ 54 29 12 35 40 91 1632 8 55 21]
```

```
[ 8 74 67 18 5 15 6 1812 4 10]
```

```
[ 42 17 26 68 24 54 155 6 1212 56]
```

```
[ 56 17 50 30 11 49 15 28 33 1306]]
```

```
[1844 5006 4279 2717 2445 2428 2031 2123 1496 1663]
```

**We can monitor the
output live**

You have learned

- how to run programs using the SLURM scheduler
- the difference between interactive and batch jobs
- how to check the status of your jobs
- how to select debug vs general SLURM partitions
- how to ask for the hardware resources you need
- you ran machine learning classifiers interactively with Jupyterhub and in batch mode

