

Intermediate Level Introduction to Computing at CARC FORTRAN

Matthew Fricke

Version 0.2




Goals

- 1) SLURM scheduler literacy
- 2) Ability to employ embarrassingly parallel solutions
- 3) Ability to employ coupled parallel solutions
- We won't cover file transfer, storage systems, module system, conda, PBS. (These are all covered in depth in the video tutorials)

Agenda

- HPC and Parallelism
- HPC Schedulers
- SLURM
- Embarrassingly Parallel
SLURM Arrays
GNU Parallel
- Coupled Parallelism
- MPI with MPI



GNUparallel



MPI

We will have one 15 minute break. Opportunity to see the machine room.

Tuesday, February 2, 20XX

Sample Footer Text

Logging into Hopper



First login to the Linux **workstation** in front of you.

Use your CARC username and password.

Jacob, Keven, Tannor, and Jose can help you login if you have trouble.

This is an “important step” so don’t let me move on until you have logged in

Logging into Hopper



```
ssh vanilla@hopper.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Logging into Hopper

Welcome to Hopper

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to <http://carc.unm.edu>, select the "New Users" menu and then click "Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the `chmod` command to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing, supported in part by the National Science Foundation, for providing the research computing resources used in this work."

Please send citations to publications@carc.unm.edu.

There are three types of slurm partitions on Hopper:
1) General - this partition is accessible by all CARC users.

2) Condo - preemptable scavenger queue available to all condo users. Your job must use checkpointing to use this queue or you will lose any work you have done if it is preempted by the partition's owner.

ENJOY

Slurmm
SODA

IT'S HIGHLY ADDICTIVE!

VOTED **#1** SOFT DRINK OF THE 31ST CENTURY!



SLURMS MCKENZIE

```
[vanilla@hopper ~]$ qgrok
```

```
queues free busy offline jobs nodes CPUs GPUs
```

```
-----  
general 1    9    0    7   10   320  0  
debug   2    0    0    0    2    64  0  
totals: 1    9    0    7   10   320  0
```



```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST  
debug      up    4:00:00    2  idle hopper[011-012]
```

sinfo reports information about
partitions

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST  
debug      up    4:00:00    2 idle hopper[011-012]
```

The debug queues are intended
for testing your programs.

And for interactive jobs.

```
[vanilla@hopper ~]$ sinfo --partition debug
```

```
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
```

```
debug      up    4:00:00    2 idle hopper[011-012]
```



Name

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST  
debug      up    4:00:00    2 idle hopper[011-012]
```



You can run a “job” for up to 4 hrs.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST  
debug      up    4:00:00    2  idle hopper[011-012]
```



There are two nodes in this partition.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST
debug      up    4:00:00    2  idle hopper[011-012]
```



The names of the nodes in the
partition


```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up    4:00:00    2  idle hopper[011-012]
```



The names of the nodes in the
partition

```
[vanilla@hopper ~]$ sinfo --partition general
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
general*	up	2-00:00:00	9	alloc	hopper[001-009]
general*	up	2-00:00:00	1	idle	hopper010



Hopper001 through 009 are running jobs.

Hopper010 is waiting to be used.

```
[vanilla@hopper ~]$ hostname  
hopper  
[vanilla@hopper ~]$
```



Running on the Head Node.

The head node's name is "hopper".

```
[vanilla@hopper ~]$ hostname
```

```
hopper
```

```
[vanilla@hopper ~]$ man hostname
```

```
[vanilla@hopper ~]$ hostname
```

```
hopper
```

```
[vanilla@hopper ~]$ man hostname
```

```
('q' to quit)
```

```
[vanilla@hopper ~]$ man man
```

```
('q' to quit)
```

```
[vanilla@hopper ~]$ man sinfo
```

sinfo(1)

Slurm Commands

sinfo(1)

NAME

sinfo - View information about Slurm nodes and partitions.

SYNOPSIS

sinfo [OPTIONS...]

DESCRIPTION

sinfo is used to view partition and node information for a system running Slurm

OPTIONS

-a, --all

Display information about all partitions. This causes information to be displayed about partitions that are configured as hidden and partitions that are unavailable to the user's group.


```
[vanilla@hopper ~]$ sinfo --all
```

```
PARTITION AVAIL TIMELIMIT  NODES  STATE NODELIST
general*   up 2-00:00:00    9 alloc hopper[001-009]
general*   up 2-00:00:00    1 idle hopper010
debug      up  4:00:00     2 idle hopper[011-012]
condo      up 2-00:00:00     1 down* hopper045
condo      up 2-00:00:00     3 mix hopper[018-020]
condo      up 2-00:00:00    16 alloc hopper[013-015,028-036,049-052]
condo      up 2-00:00:00    18 idle hopper[016-017,021-027,037-044,053]
bugs       up 7-00:00:00     2 alloc hopper[013-014]
pcnc       up 7-00:00:00     1 alloc hopper015
pcnc       up 7-00:00:00     1 idle hopper016
pathogen   up 7-00:00:00     1 idle hopper017
tc         up 7-00:00:00     3 mix hopper[018-020]
tc         up 7-00:00:00     2 alloc hopper[029-030]
tc         up 7-00:00:00     5 idle hopper[021-025]
gold       up 7-00:00:00     2 idle hopper[026-027]
fishgen    up 7-00:00:00     1 alloc hopper028
neuro-hsc  up 7-00:00:00     6 alloc hopper[031-036]
neuro-hsc  up 7-00:00:00     8 idle hopper[037-044]
cup-ecs    up 7-00:00:00     2 alloc hopper[049-050]
tid        up 7-00:00:00     1 alloc hopper051
biocomp    up 7-00:00:00     1 alloc hopper052
chakra     up 7-00:00:00     1 idle hopper053
pna        up 7-00:00:00     1 down* hopper045
```

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Tell slurm to run a program
on a compute node...

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Run the program on a
compute node in the
debug partition.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



The program
to run.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
[vanilla@hopper ~]$ queue
```



```
[vanilla@hopper ~]$ queue
```

JOBID	PARTITION	NAME	USER	ST	TIME
4314	general	PRE	erowland	PD	0:00
4315	general	PRE	erowland	PD	0:00
4317	general	PRE	erowland	PD	0:00
4318	general	PRE	erowland	PD	0:00

**PD means programs
that are waiting their
turn.**

Shows you what the slurm scheduler is doing right now.

Here we can see that user 'erowland' has a lot of programs waiting to run.

[illegible]

```
[vanilla@hopper ~]$ queue
```

[illegible]

```
[vanilla@hopper ~]$ squeue -t R --all
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
4405	condo	2ndMA	mfricke	R	1-07:48:30	6	hopper[031-036]
5208	condo	NN	kgu	R	5:48:49	1	hopper015
5210	condo	NN	kgu	R	6:30:13	1	hopper014
5209	condo	NN	kgu	R	6:31:13	1	hopper013
5206	condo	NN	kgu	R	6:32:13	1	hopper051
5207	condo	NN	kgu	R	6:32:13	1	hopper052
5205	condo	NN	kgu	R	6:32:43	1	hopper028
4595	cup-ecs	golConfi	aalasand	R	2-06:51:59	1	hopper050
4594	cup-ecs	golConfi	aalasand	R	2-06:52:03	1	hopper049
5120	general	jupyterh	jacobm	R	11:45:47	1	hopper007
4313	general	PRE	erowland	R	1:17:29	2	hopper[003-004]
5111	general	1stMA	mfricke	R	11:15:28	2	hopper[005-006]
5025	general	c2n	jxzuo	R	1:50	1	hopper001
5024	general	c2n	jxzuo	R	31:28	1	hopper002
5203	general	NN	kgu	R	6:37:50	1	hopper009
5201	general	NN	kgu	R	6:38:14	1	hopper008
4390	tc	UCsTpCyd	lepluart	R	2-15:18:18	3	hopper[018-020]
5198	tc	NN	kgu	R	6:40:19	1	hopper030
5196	tc	NN	kgu	R	6:40:31	1	hopper029

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

You ran two **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

You ran eight **copies** of your program.

ntasks is the number of copies to run.


```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
```

```
srun: Account not specified in script or ~/.default_slurm_account, using latest project
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

By default, each task (copy of your program) is allowed to use one CPU.

Many programs are able to use more than one CPU at a time.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.  
hopper011  
hopper011
```

**Here we are telling SLURM to
run 2 copies of our program
and let each copy of our
program use 2 CPUs.**

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2 --ntasks-per-node 4 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
hopper012
```

You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.

```
hopper012
```

```
hopper011
```

```
hopper011
```

```
hopper012
```

```
hopper012
```

```
hopper011
```

```
hopper011
```

**Here we are telling SLURM to run 4
copies of our program on 2 different
compute nodes.**

**This is useful when our programs need
a bigger share of the compute node.**

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2  
--ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper011
```

You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.

```
hopper011
```

```
hopper012
```

```
hopper012
```

And we can combine all three.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been allocated  
submission script.
```

```
hopper011
```

```
Hopper011
```

**And we can specify how much
memory we want.**

**--mem 4G means give me 4
gigabytes of memory per node.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been allocated  
submission script.
```

```
hopper011
```

```
Hopper011
```

Why does all this matter?

**The purpose of SLURM is to provide
you the hardware your programs
need.**

**So you have to understand what
those requirements are really well.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project
```

```
hopper012
```

```
hopper012
```

```
You have not been allocated  
submission script.
```

```
hopper011
```

```
Hopper011
```

- 1) Can my program use multiple CPUs?
- 2) How much memory does my program need?
- 3) Can my program use multiple compute nodes (MPI*, GNU Parallel*)?
- 4) Can my program use GPUs?

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest  
project  
hopper012  
hopper012  
You have not been allocated a  
submission script.  
hopper011  
Hopper011
```

This command is getting pretty long.

We can use **salloc** to avoid asking for the same resources every time we use **srun**.


```
[vanilla@hopper ~]$ salloc --partition debug --nodes 2 --ntasks-per-node 2  
salloc: Account not specified in script or ~/.default_slurm_account, using latest  
project  
salloc: Granted job allocation 5251  
salloc: Waiting for resource configuration  
salloc: Nodes hopper[011-012] are ready for job  
[vanilla@hopper ~]$
```

This command is getting pretty long.

We can use **salloc** to avoid asking for the same resources every time we use **srun**.

```
[vanilla@hopper ~]$ srun hostname
```

```
hopper012
```

```
hopper012
```

```
hopper011
```

You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.

```
hopper011
```

```
[vanilla@hopper ~]$ srun hostname
```

```
hopper012
```

```
hopper011
```

```
hopper012
```

```
hopper011
```

```
[vanilla@hopper ~]$
```

Now we can use **srun** over and over without having to ask for a new hardware allocation each time.

```
[vanilla@hopper ~]$ exit
```

```
exit
```

```
salloc: Relinquishing job allocation 5251
```

Always type **exit** when you are done with the hardware.

Running `salloc` inside an allocation gets very confusing.

Interactive vs Batch Mode

Interactive Mode

- Everything so far has been interactive. You request hardware, run your program, and get the output on your screen right away.

Batch Mode

- Most programs at an HPC center are run in “batch” mode.
- Batch mode means we write a shell script that the SLURM scheduler runs for us. The script requests hardware just like we did with `salloc` and then runs the commands in the script.
- Whatever would have been written to the screen is saved to a file instead.

```
[vanilla@hopper ~]$ git clone https://lobogit.unm.edu/CARC/workshops.git
```

```
Cloning into 'workshops'...
```

```
remote: Enumerating objects: 132, done.
```

```
remote: Counting objects: 100% (75/75), done.
```

```
remote: Compressing objects: 100% (43/43), done.
```

```
remote: Total 132 (delta 33), reused 74 (delta 32), pack-reused 57
```

```
Receiving objects: 100% (132/132), 57.58 KiB | 3.60 MiB/s, done.
```

```
Resolving deltas: 100% (51/51), done.
```

Rather than make you write shell scripts lets just download some we wrote for this workshop...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H2O.gjf
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run tree to see how the workshops directories are organized...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H2O.gjf
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run **tree** to see how the workshops directories are organized...

The workshop files are divided into “code”, “slurm”, and “data” directories.

```
[vanilla@hopper intro_workshop]$ pwd
/users/vanilla/workshops/intro_workshop
[vanilla@hopper intro_workshop]$ cat slurm/workshop_example1.sh
#!/bin/bash
#SBATCH --partition debug
#SBATCH --ntasks 4
#SBATCH --time 00:05:00
#SBATCH --job-name ws_example
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

srun hostname
```

Let's take a look at the **workshop_example.sh** script in the slurm directory...


```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh
sbatch: Account not specified in script or ~/.default_slurm_account, using latest project
Submitted batch job 5252
[vanilla@hopper intro_workshop]$
```

We **submit** our slurm
shell script with the
sbatch command.

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh  
sbatch: Account not specified in script or ~/.default_slurm_account, using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

Notice that the only output we get is a job id.

This indicates that the script was successfully sent to the scheduler.

The commands in the script will run as soon as the hardware requested is available.

We **submit** our slurm shell script with the sbatch command.

Workflow

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Compute Node 01

Compute Node 02

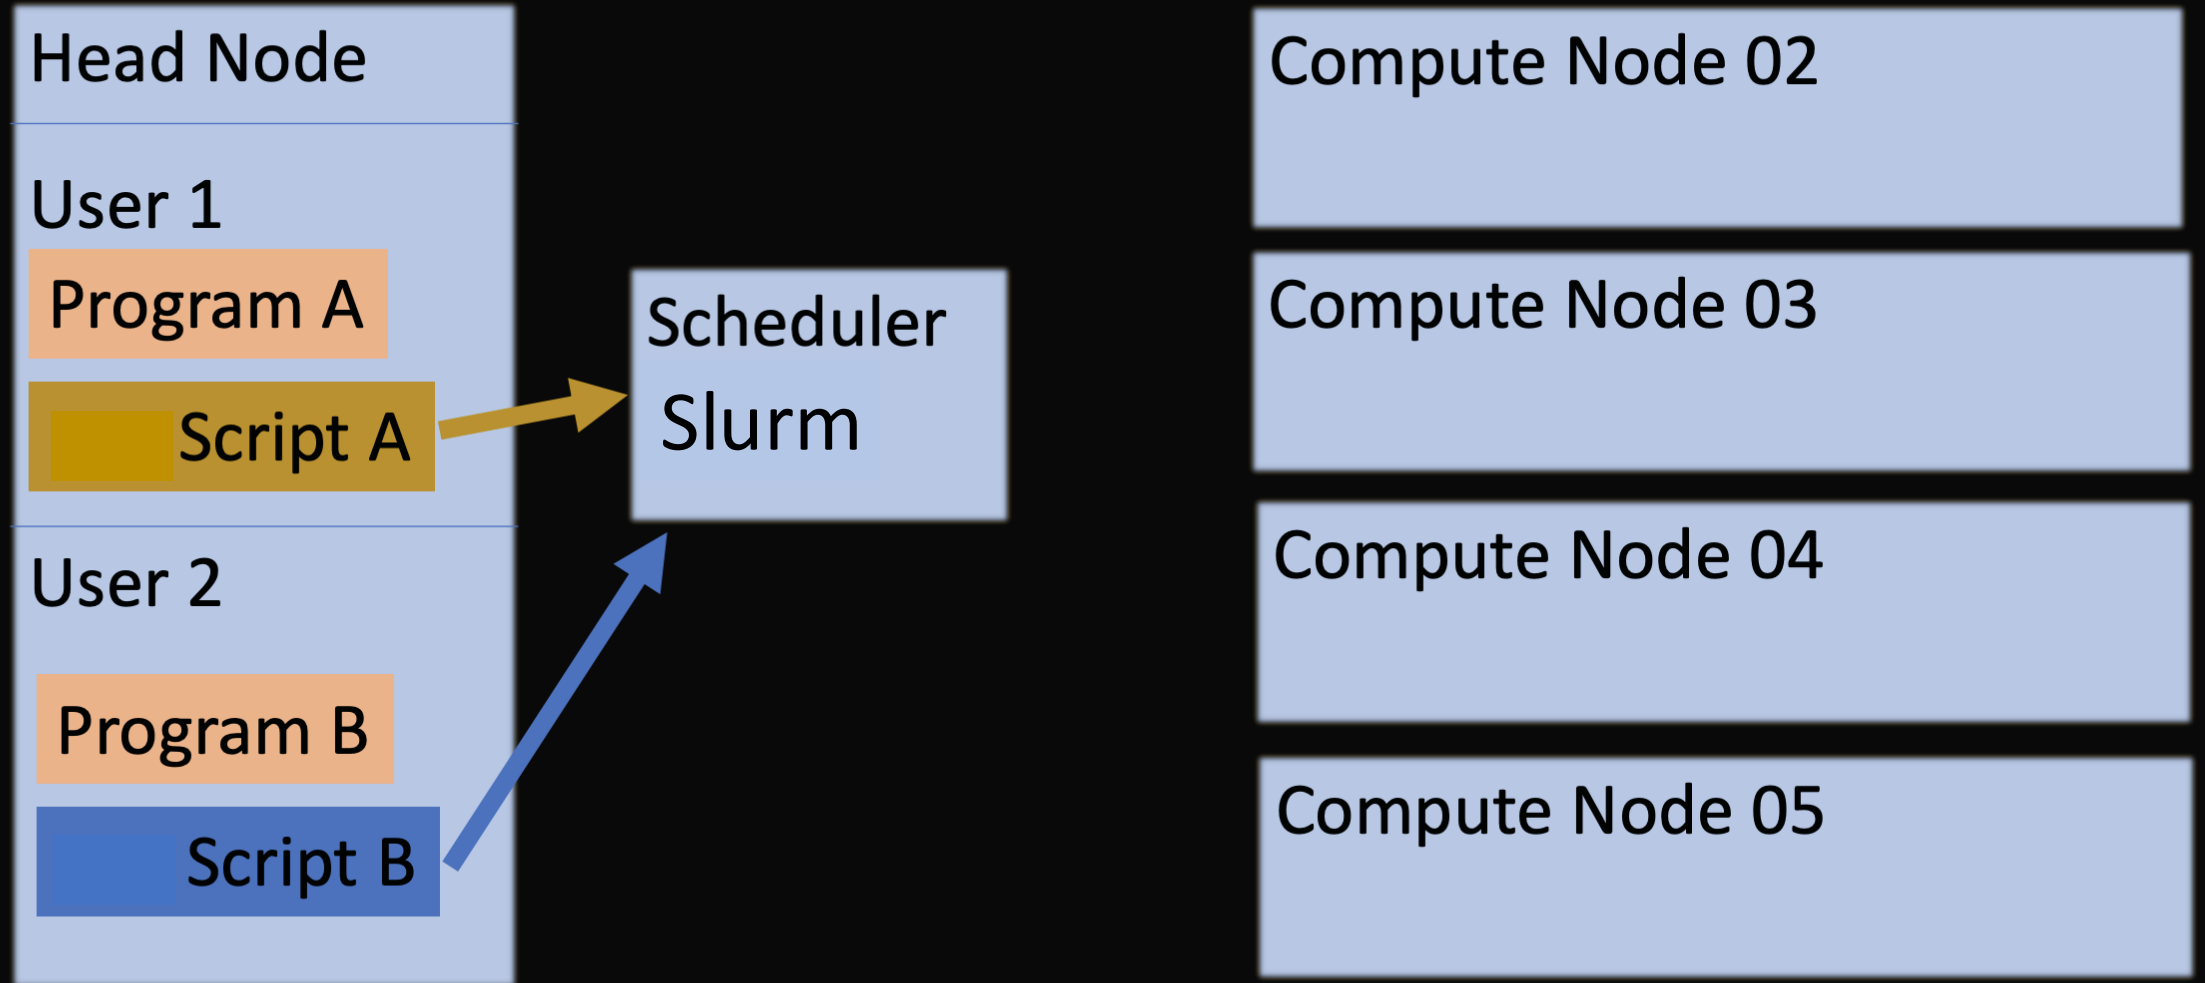
Compute Node 03

Compute Node 04

Compute Node 05

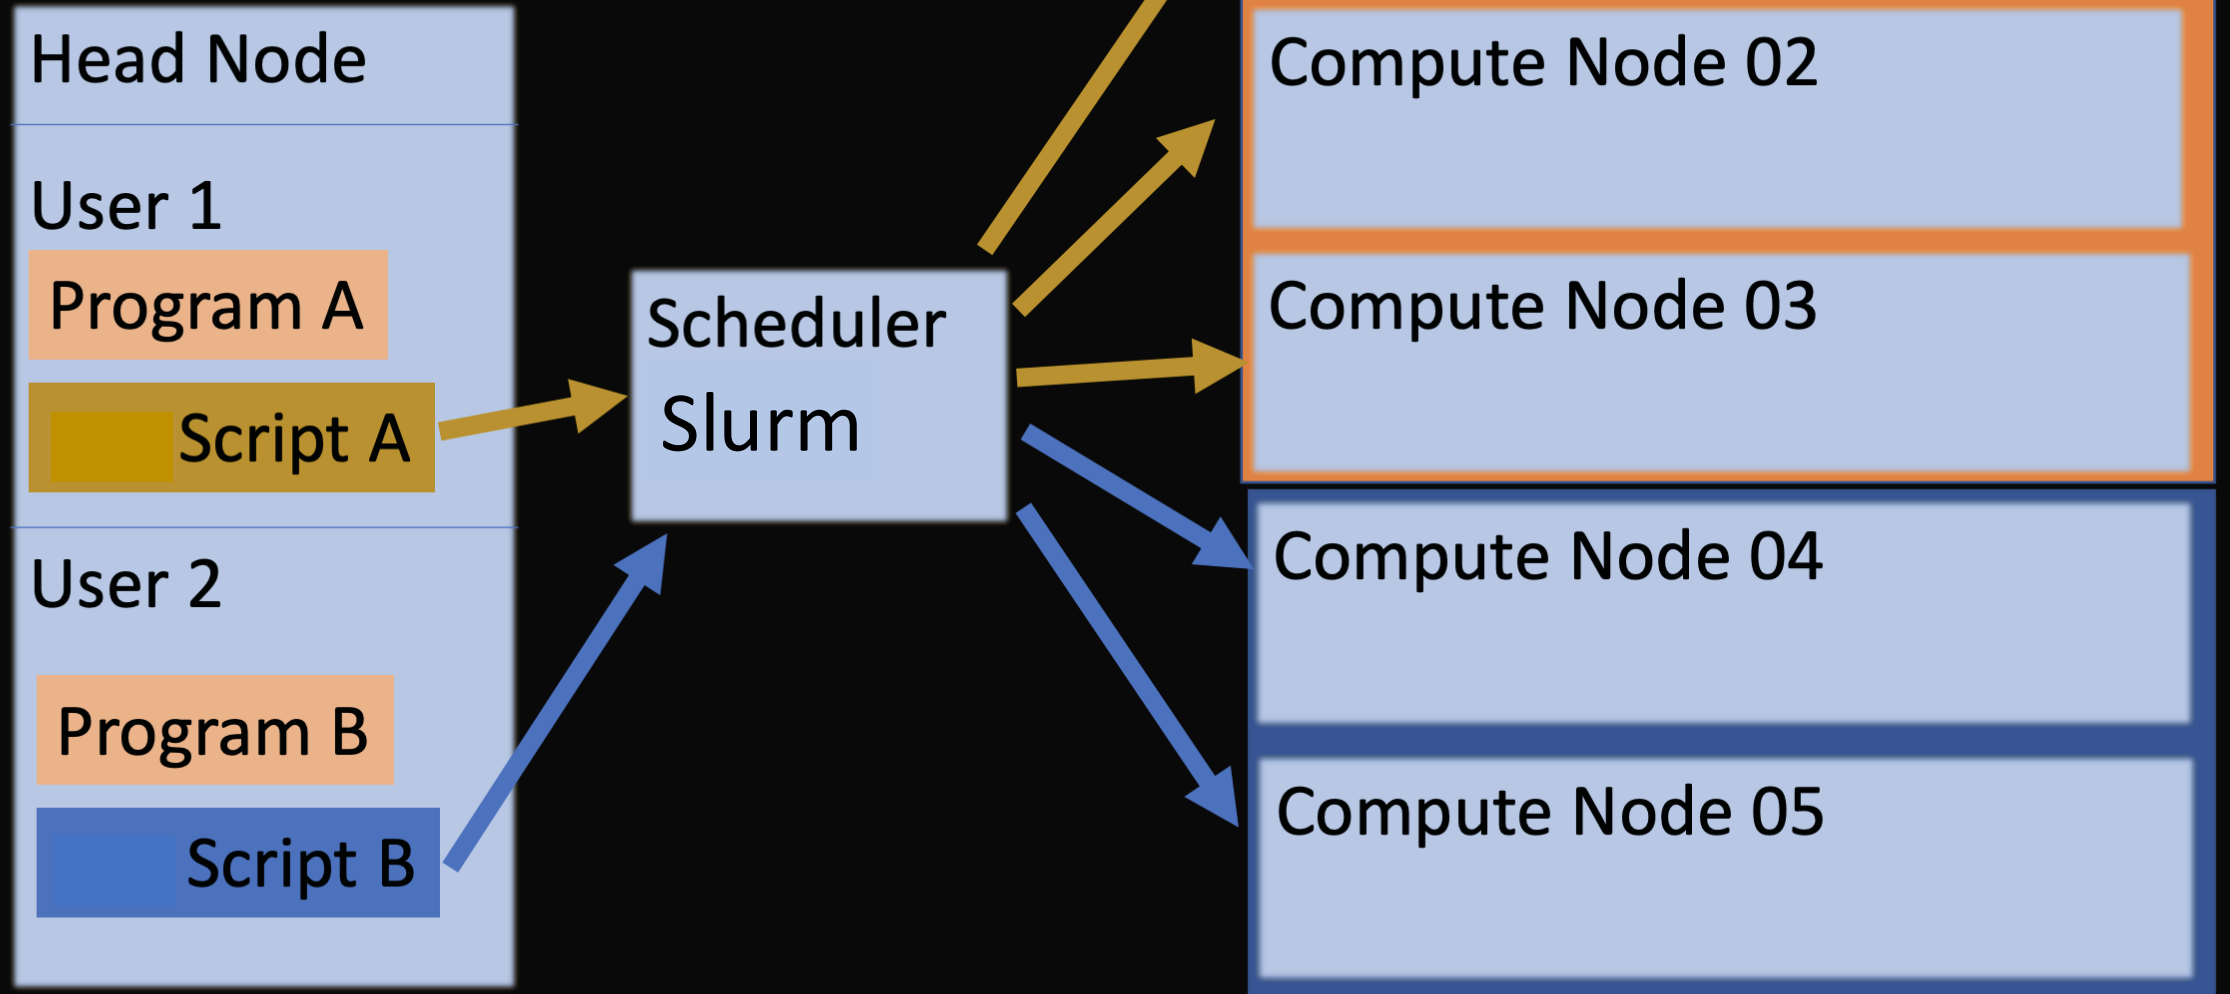
Shared filesystems – All nodes can access the same programs and write output

Workflow



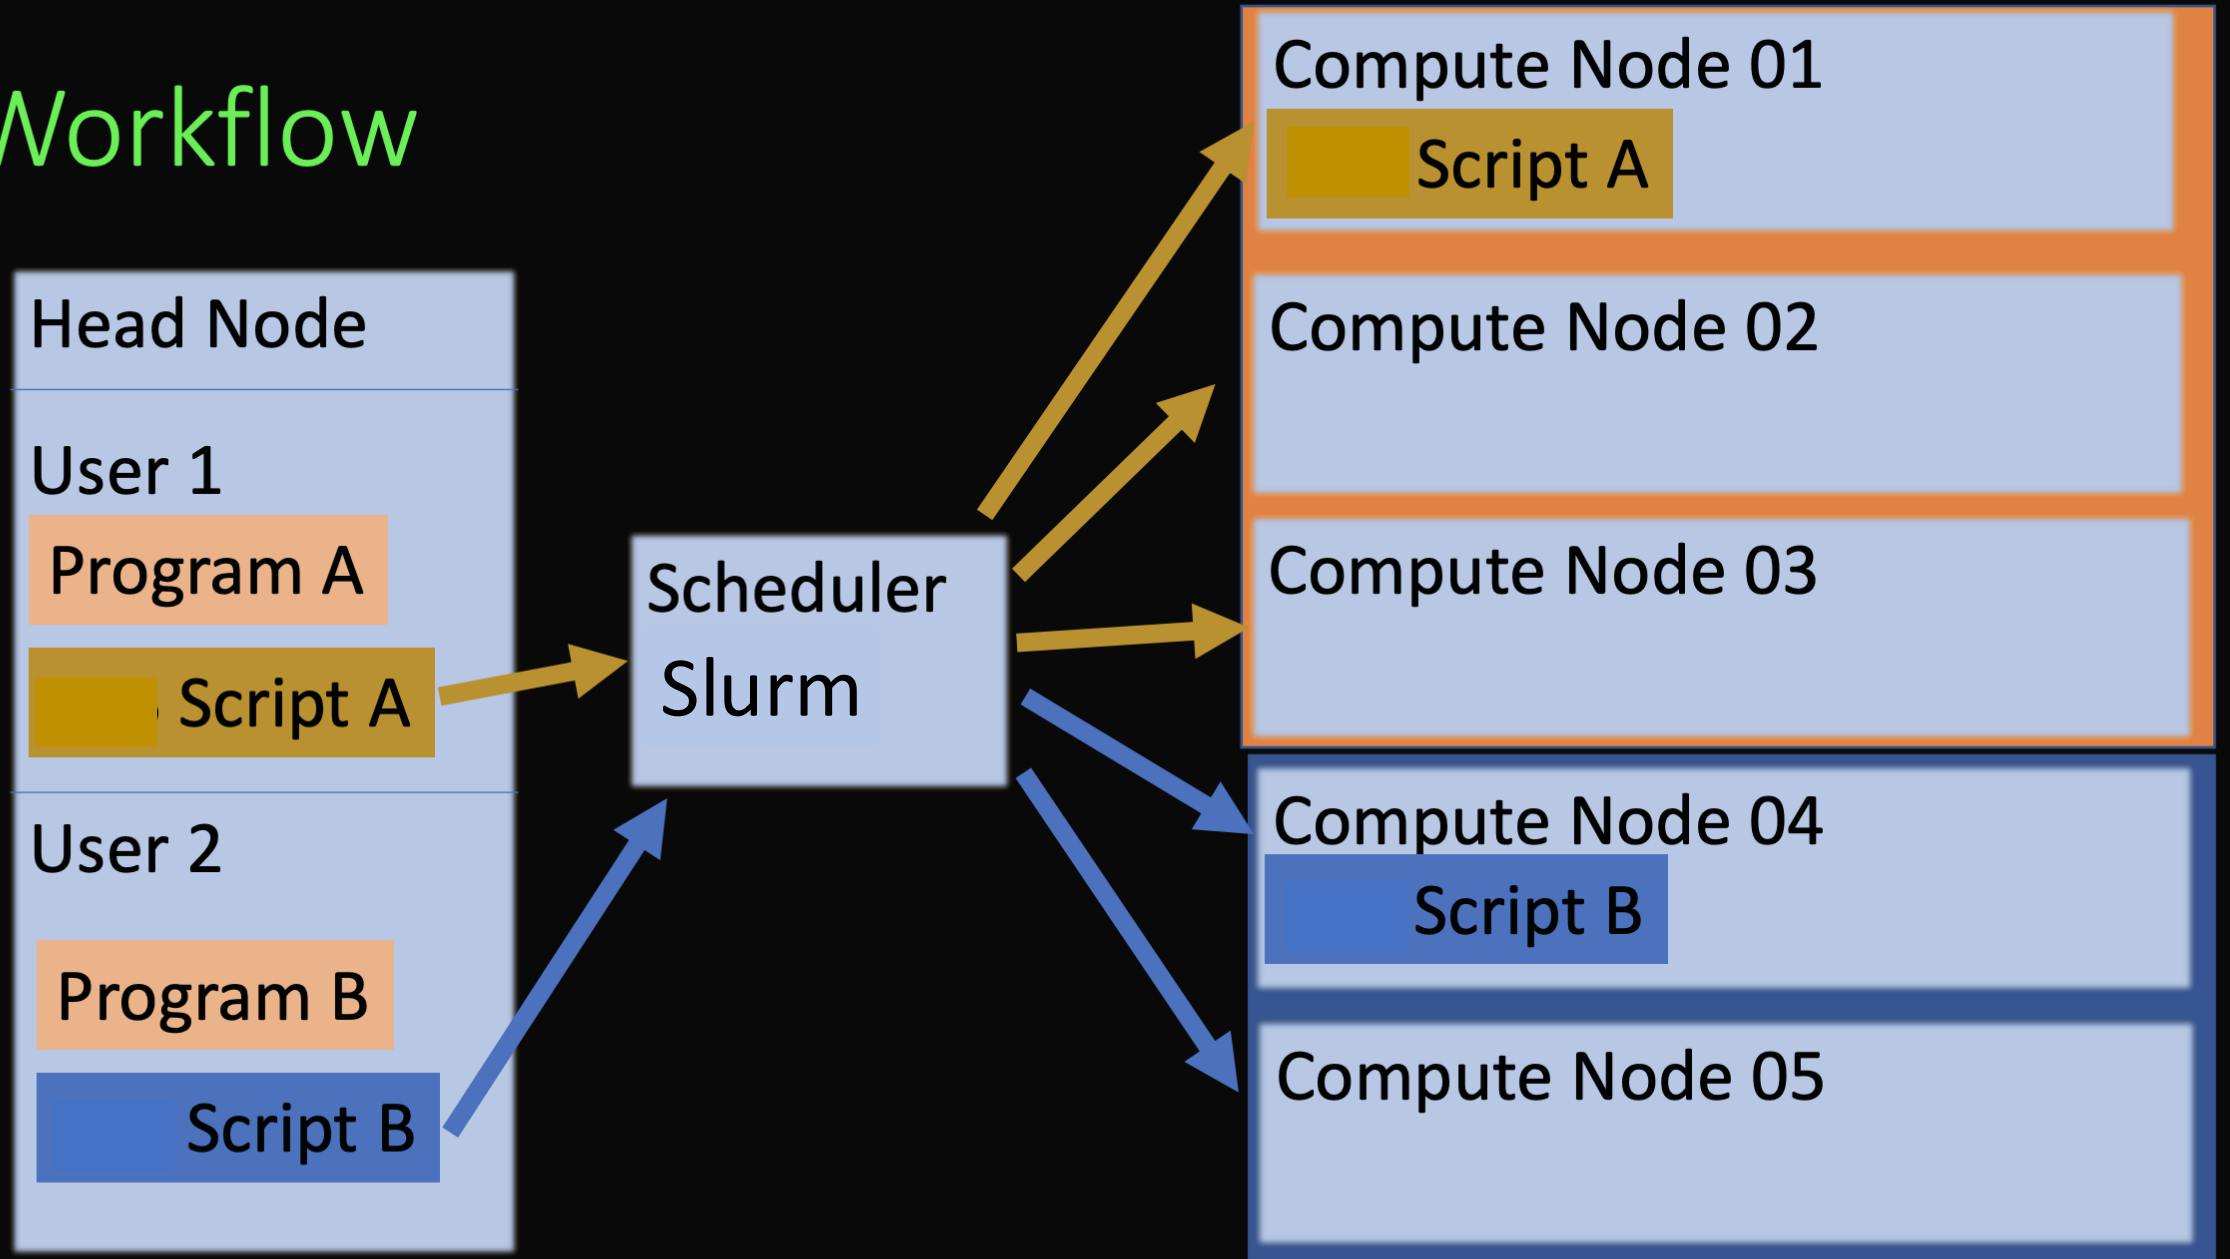
Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

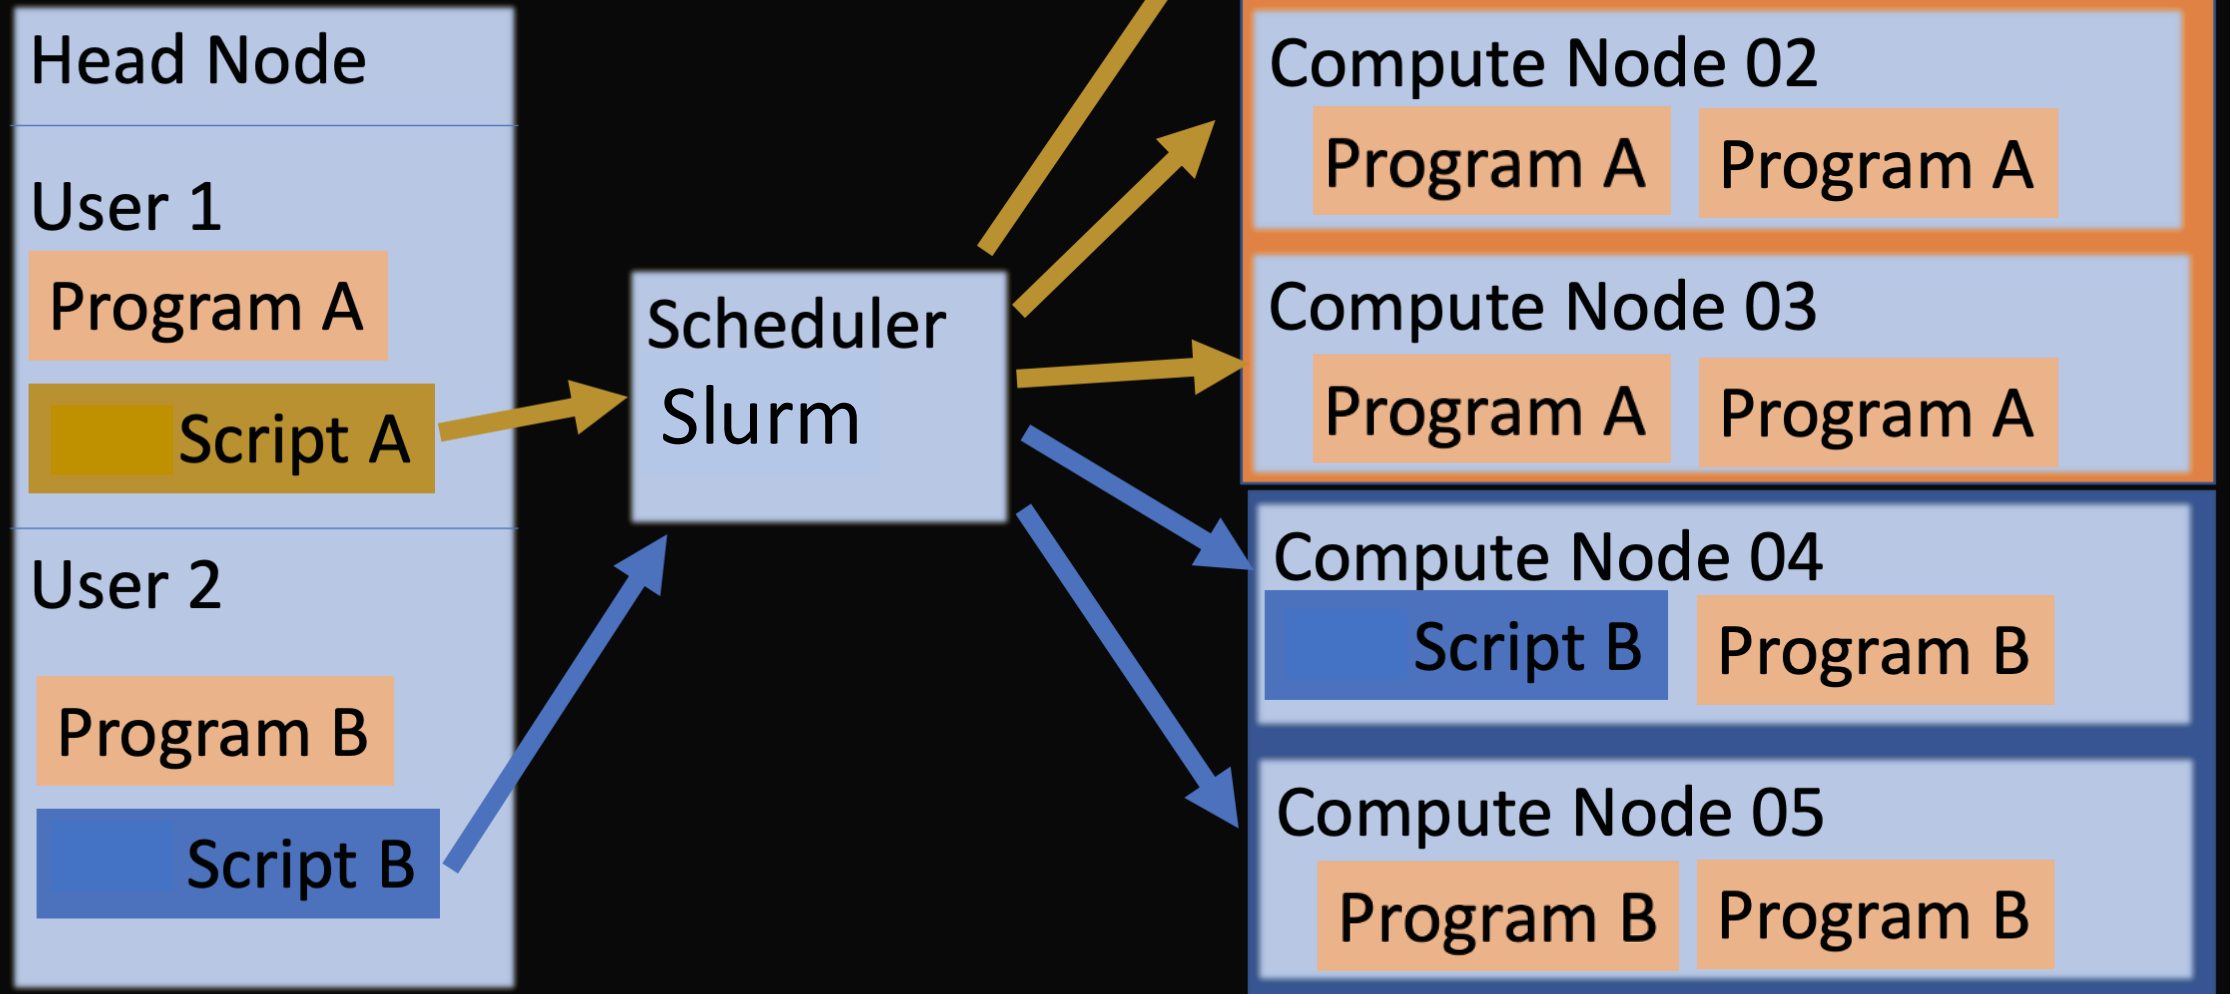
Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow


We need something in the script to run the program on all the nodes. E.g. `srun`.



Shared filesystems – All nodes can access the same programs and write output

```
[vanilla@hopper intro_workshop]$ ls  
code data pbs slurm slurm-5252.out
```

The **hostname**
command is very fast
so everyone's job
should finish in a few
seconds.



When it is finished you will
have a new file named
slurm-{your job id}.out.


```
[vanilla@hopper intro_workshop]$ ls  
code data pbs slurm slurm-5252.out
```



When it is finished you will have a new file named `slurm-{your job id}.out`.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
[vanilla@hopper intro_workshop]$ ls  
code data pbs slurm slurm-5252.out
```




When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
What do you see?
```

```
[vanilla@hopper intro_workshop]$ cd code  
[vanilla@hopper code]$ pwd  
~/workshops/intro_workshop/code
```

**Let's experiment with a
program that does
slightly more than
print the hostname.**

Serial Program to Calculate π



$$\int_0^1 \frac{4}{1+x^2} = \pi^*$$

$$\sum_{i=1}^N \frac{4}{1 + \left(\frac{i+\frac{1}{2}}{N}\right)^2} \approx \pi$$

*Because arctan...

Program to estimate π – *Main program*

```
[vanilla@hopper code]$ cat calc_pi_serial.f90
```

```
program cal_pi_serial
```

```
  implicit none
```

```
  integer steps
```

```
  character(len=100) :: steps_arg
```

```
  real :: start, finish
```

```
  double precision :: p, pi_ref = 3.14159265358979323846264338
```

```
  call getarg(1, steps_arg) ! Read argument
```

```
  read(steps_arg,*) steps ! Convert string to integer
```

```
  call cpu_time(start) ! Time how long it takes to est pi
```

```
  p = pi(steps)
```

```
  call cpu_time(finish)
```

```
  print '("Pi=", g0, " (Diff=", g0,"))', p, abs(p-pi_ref)
```

```
  print '("calculated in ", g0,"s with ", g0, " steps)",
```

```
    finish-start, steps
```

```
contains
```

Program to estimate π – *Function pi()*

! The area under the curve $4/(1+x^2)$ is π .

! Estimate with the Riemann sum.

```
function pi(num_steps) result(area)
double precision x, area, step, sum
integer num_steps
sum = 0
step_size = 1.0/num_steps ! Determine the rectangle size

! Loop over the rectangles under the curve
do i = 0, num_steps
    x = (i + 0.5)*step_size ! Calc the ith rectangle area
    sum = sum + 4.0/(1.0+x*x) ! Sum them up
end do

area = step_size*sum ! Normalise to be between 0 and 1

end function

end program
```

Program to estimate π – *Function pi()*

! The area under the curve $4/(1+x^2)$ is π .

! Estimate with the Riemann sum.

```
function pi(num_steps) result(area)
```

```
double precision x, area, step, sum
```

```
integer num_steps
```

```
sum = 0
```

```
step_size = 1.0/num_steps ! Determine the rectangle size
```

! Loop over the rectangles under the curve

```
do i = 0, num_steps
```

```
  x = (i + 0.5)*step_size ! Calc the ith rectangle area
```

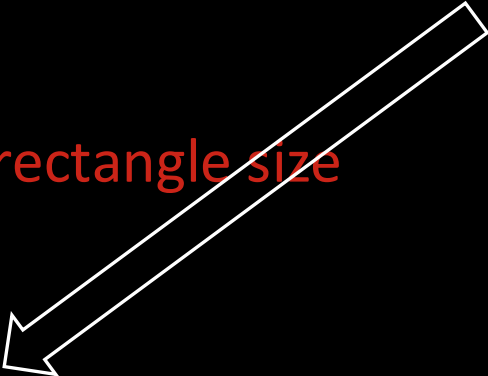
```
  sum = sum + 4.0/(1.0+x*x) ! Sum them up
```

```
end do
```

```
area = step_size*sum ! Normalise to be between 0 and 1
```

```
end function
```

```
end program
```

$$\sum_{i=1}^N \frac{4}{1 + \left(\frac{i + \frac{1}{2}}{N}\right)^2} \approx \pi$$


```
[vanilla@hopper code]$ module load intel/20.0.4
```

```
[vanilla@hopper code]$ ifort calc_pi_serial.f90 -o calc_pi_serial
```

```
[vanilla@hopper code]$ srun --partition debug ./calc_pi_serial 1000
```

```
srun: Using account 2016199 from ~/.default_slurm_account
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.
```

```
Pi=3.143591832167984 (Diff=.1999091155410415E-02)
```

```
(calculated in .2999790E-05s with 1000 steps)
```


Load intel compilers then compile our FORTRAN program.

Run calcPiSerial.py on a compute node.

For our example program the more steps it takes the more accurate it is, but the longer it takes.


Parallelism – Embarrassingly Parallel

- Embarrassingly parallel (Cleve Moler) are problems that are really really easy to speed up with more CPUs.
- The most common example is that you have a program that runs in serial and takes some input file, processes it, and produces some output.
- The problem is that you have 1,000 of the input files and want to run your program on each one.



Parallelism – Embarrassingly Parallel

This is “embarrassing”
because all you have to do
is run 1,000 copies of your
program on 1,000 CPUs
each with a different input
file and you are done.



SLURM ARRAYS

- One way to run the 1,000 copies of your program on 1,000 different inputs would be to write 1,000 slurm scripts each specifying a different input to your program and then sbatch submit them all. (this would work but there are better ways).
- SLURM arrays are used to schedule a lot of jobs with one slurm script.

```
[vanilla@hopper intro_workshop]$ nano slurm/calc_pi_array_f90.sh
```

```
#!/bin/bash
```

```
#SBATCH --partition debug
```

```
#SBATCH --ntasks 1
```

```
#SBATCH --time 00:05:00
```

```
#SBATCH --job-name calc_pi_array
```

```
#SBATCH --mail-user your_username@unm.edu
```

```
#SBATCH --mail-type ALL
```

```
#SBATCH --array=1-12%12
```

```
echo "$HOSTNAME - $SLURM_ARRAY_TASK_ID"
```

```
module load intel/20.0.4 intel-mpi/2019.10.317-ruxn
```

```
NUM_STEPS="${SLURM_ARRAY_TASK_ID}0000"
```

```
echo "Calculating pi with $NUM_STEPS..."
```

```
code/calc_pi_serial $NUM_STEPS
```

Requires some annoying bash scripting.

\$something means get the value of the variable "something"

--array says

- 1) run 12 separate jobs**
- 2) Store the count of the job in the variable SLURM_ARRAY_TASK_ID**

```
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh
```

```
sbatch: Using account 2016199 from ~/.default_slurm_account
```

```
Submitted batch job 5263
```

```
vanilla@hopper:~/workshops/intro_workshop$ squeue -me
```

```
JOBID PARTITION  NAME  USER ST  TIME  NODES NODELIST(REASON)
```

```
5263  debug calc_pi_ vanilla R   0:44   1 hopper011
```

Submit the array script.

Then enter **squeue --me** to see the job status.

Take a look at the job output. (How many output files do you have?)

JOB arrays are OK or very simple inputs like programs that take a single file as input. But even passing in a value takes some annoying variable manipulation.


GNU Parallel is much more sophisticated it can take inputs in all sorts of ways. We will look at just 3 ways.

To access GNU parallel enter **module load parallel**

Let's experiment with parallel interactively...

***NOTE: we don't use srun to run parallel.**

GNU Parallel



Has been around for a very long time and has lots and lots of great features.

But basically it creates a job for every input it receives. The inputs can be specified in the command, read from a file, or be the output of another program.

It also remembers which jobs have finished and which still need to be run. So when you run out of time and resubmit it will automatically pick up where it left off.


```
[vanilla@hopper intro_workshop]$ seq 10 10 100
```

```
10  
20  
30  
40  
50  
60  
70  
80  
90  
100
```

GNU Parallel can read the output of other programs and use them as inputs to your program.

Here a copy of calc pi is run for each row in the output of **seq**

```
[vanilla@hopper intro_workshop]$ module load parallel
```

```
[vanilla@hopper intro_workshop]$ seq 10 10 100 | parallel code/calc_pi_serial
```

```
Pi = 3.14180098689309428295, (Diff=0.00020833330330116695) (calculated in 0.000007 secs with 20 steps)
```

```
Pi = 3.14242598500109870940, (Diff=0.00083333141130559341) (calculated in 0.000006 secs with 10 steps)
```

```
Pi = 3.14168524617974842528, (Diff=0.00009259258995530928) (calculated in 0.000007 secs with 30 steps)
```

```
etc
```



```
[vanilla@hopper intro_workshop]$ find -name *.sh
```

```
./slurm/calc_pi_array.sh  
./slurm/calc_pi_mpi.sh  
./slurm/calc_pi_parallel.sh  
./slurm/calc_pi_serial.sh  
./slurm/gaussian.sh  
./slurm/hostname_mpi.sh  
etc
```

```
$ find -name *.sh | parallel wc -l  
7 ./code/vecadd/vecaddmpi_cpu.sh  
19 ./slurm/calc_pi_array.sh  
15 ./slurm/calc_pi_mpi.sh  
20 ./slurm/calc_pi_parallel.sh  
14 ./slurm/calc_pi_serial.sh  
16 ./slurm/gaussian.sh  
15 ./slurm/hostname_mpi.sh  
etc
```

A common application is to use **find** to produce a list of paths with some extension.

Then parallel runs some program on each path.

In this case **wc -l** counts the lines in a file. In some real CARC examples the input files are phylogenetic trees, graphs, neuroimages, or CT scans.

```
[vanilla@hopper intro_workshop]$ exit  
exit  
salloc: Relinquishing job allocation 5275
```

**Don't forget to exit your
salloc allocation.**

```
$sbatch slurm/calc_pi_parallel_f90.sh
```

```
sbatch: Using account 2016199 from ~/.default_slurm_account
```

```
Submitted batch job 5276
```

Submit the job, check it's progress with `squeue --me`, and take a look at the output.

```
$sbatch slurm/calc_pi_parallel_f90.sh
```

```
sbatch: Using account 2016199 from ~/.default_slurm_account
```

```
Submitted batch job 5276
```


**Submit the job, check it's progress with `squeue --me`,
and take a look at the output.**

Parallelism – Coupled Parallelism


- Coupled problems are those where the CPUs need to work together to solve a problem by communicating with each other.
- Many commercial and research programs designed to run on HPC systems like CARC use a library called the message passing interface (MPI) to do this.
- We have written an MPI version of our python pi calculator to demonstrate.



Serial Program to Calculate π



Parallel Program to Calculate π



MPI Program to estimate π

```
program cal_pi_mpi
  use mpi ! Import Message Passing Interface
  implicit none ! Don't allow FORTRAN to infer types
  integer steps, ierr, num_procs, root, rank
  character(len=100) steps_arg
  real :: start, finish
  double precision :: p, pi_ref = 3.141592653589793238462643383279502884197

  ! Setup Message Passing Interface
  call MPI_Init( ierr )
  call MPI_Comm_rank( MPI_COMM_WORLD, rank, ierr )
  call MPI_Comm_size( MPI_COMM_WORLD, num_procs, ierr )
  root = 0 ! Call the 0th process "root"

  call getarg(1, steps_arg) ! Read argument
  read(steps_arg,*) steps ! Convert string to integer

  call cpu_time(start) ! Time how long it takes to est pi
  ! Call the pi function and print the result
  p = pi(steps, num_procs, rank)
  call cpu_time(finish)

  if (rank==root) then
    print '("Pi=", g0, " (Diff=", g0,")"', p, abs(p-pi_ref)
    print '("(calculated in ", g0,"s with ", g0, " steps and ", g0, " processes)"', finish-start, steps, num_procs
  endif
```

```
contains
  ! The area under the curve 4/(1+x^2) is pi
  ! Estimate with the Riemann sum
  function pi(num_steps, num_procs, rank) result(area)
    double precision x, area, my_area, step, sum, step_size
    integer num_steps, num_procs, rank, i

    sum = 0
    step_size = 1.0/num_steps ! Normalise the area between 0 and 1
    do i = rank, num_steps, num_procs ! Only loop over our part of the sum
      x = (i + 0.5)*step_size
      sum = sum + 4.0/(1.0+x*x)
    end do
    my_area = step_size*sum

    ! Get that partial sums from all the processes, add them up, and give to the root
    ! syntax: input variable, output variable, data size, MPI Type, Operation, output rank, MPI World, error int
    call mpi_reduce(my_area,area,1,MPI_DOUBLE,MPI_SUM,root,MPI_COMM_WORLD,ierr)

  end function
end program
```



```
[vanilla@hopper code]$ module load intel-mpi/2019.10.317-ruxn  
[vanilla@hopper code]$ mpiifort calc_pi_mpi.f90 -o calc_pi_mpi
```

Compile the code with Intel MPI FORTRAN compiler

MPI: Message Passing Interface

When programs need to run on many processors but also communicate with one another.

Here the parallel version of calcPi needs to communicate the partial sums computed by each process so they can all be added up.

To communicate we will use the MPI library:

```
module load minconda3  
conda create -n mpi_numpy mpi mpi4py numpy
```

```
#!/bin/bash
```

```
#SBATCH --partition debug
```

```
#SBATCH --nodes 2
```

```
#SBATCH --ntasks-per-node 4
```

```
#SBATCH --time 00:05:00
```

```
#SBATCH --job-name calc_pi_mpi
```

```
#SBATCH --mail-user your_username@unm.edu
```

```
#SBATCH --mail-type ALL
```

```
module load intel/20.0.4 intel-mpi/2019.10.317-ruxn
```

```
srun --mpi=pmi2 code/calc_pi_mpi 1000000000
```

sbatch slurm/calc_pi_mpi_f90.sh

```
[vanilla@hopper intro_workshop]$ module load parallel
```

```
[vanilla@hopper code]$ seq 1 1 6 | parallel srun --partition debug --ntasks {} calc_pi_mpi 100000
```

```
Job 3056979 running on hopper011
```

```
Pi=3.141612602973879 (Diff=.1986196130587814E-04)
```

```
(calculated in .2641998E-02s with 100000 steps and 2 processes)
```

```
Job 3056977 running on hopper011
```

```
Pi=3.141612602973873 (Diff=.1986196130010498E-04)
```

```
(calculated in .2910048E-03s with 100000 steps and 1 processes)
```

```
Job 3056976 running on hopper011
```

```
Pi=3.141612602973876 (Diff=.1986196130232543E-04)
```

```
(calculated in .4189983E-03s with 100000 steps and 3 processes)
```

```
Job 3056978 running on hopper011
```

```
Pi=3.141612602973876 (Diff=.1986196130321360E-04)
```

```
(calculated in .2626002E-02s with 100000 steps and 6 processes)
```

```
Job 3056980 running on hopper011
```

```
Pi=3.141612602973874 (Diff=.1986196130054907E-04)
```

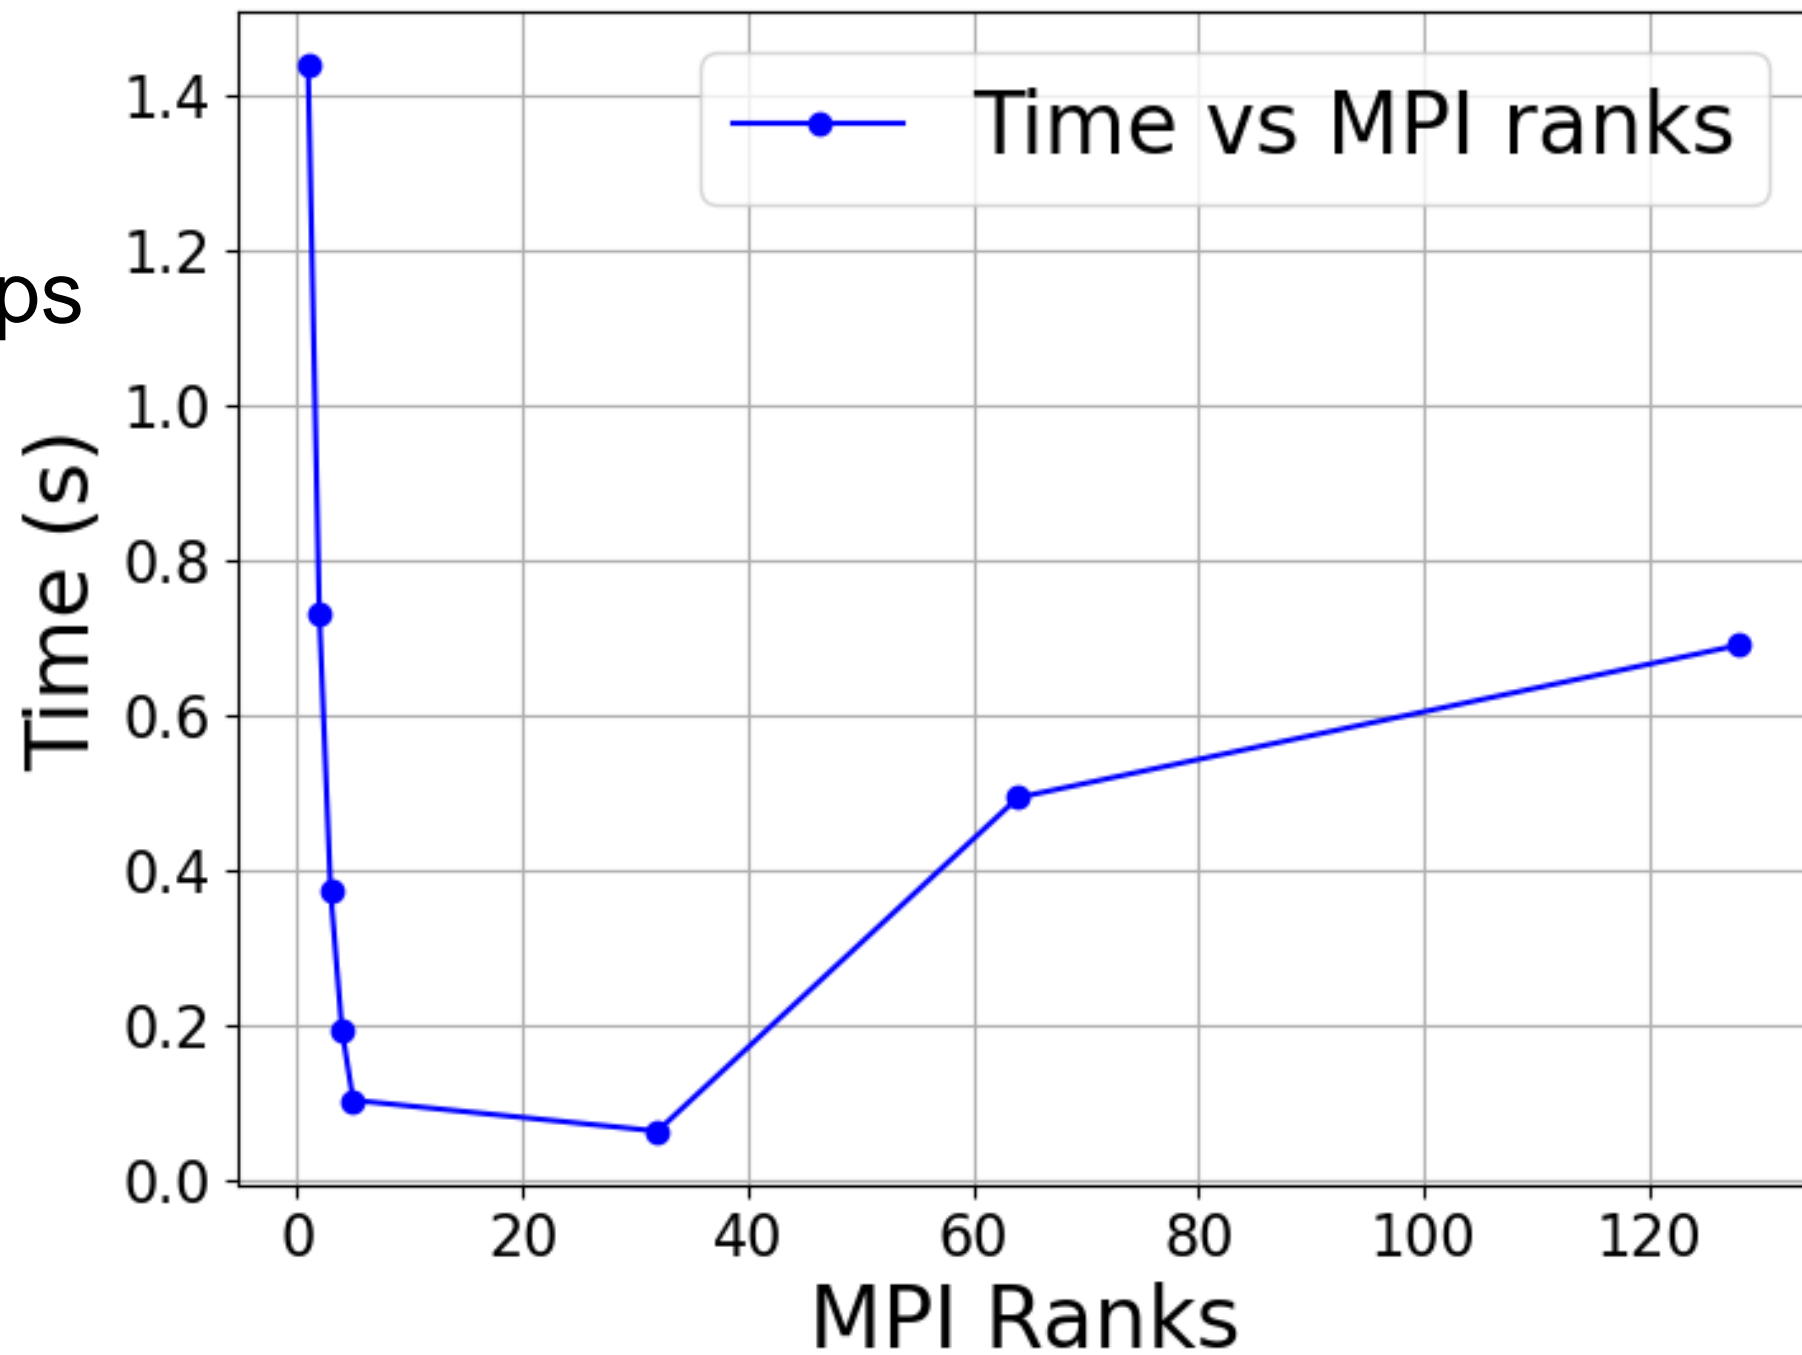
```
(calculated in .3805004E-02s with 100000 steps and 4 processes)
```

```
Job 3056981 running on hopper011
```

```
Pi=3.141612602973876 (Diff=.1986196130232543E-04)
```

```
(calculated in .2314001E-02s with 100000 steps and 5 processes)
```

Speedup with
MPI Rank
for 2×10^{10} steps



Useful Slurm Commands

<code>squeue --me --long</code>	shows information about jobs you submitted
<code>squeue --me --start</code>	shows when slurm expects your job to start
<code>scancel jobid</code>	Cancels a job
<code>scancel --u \$USER</code>	Cancels all your jobs
<code>sacct</code>	Shows your job history
<code>seff jobid</code>	Shows how efficiently the hardware was used