

Intermediate Level Introduction to Computing at CARC

Second in a three-part series – 1hr version

Matthew Fricke

Version 0.1b

b



Goals

- 1) SLURM scheduler literacy
- 2) Message Passing Interface (MPI)
- We won't cover file transfer, storage systems, module system, conda, PBS. (These are all covered in depth in the video tutorials)



Logging into Hopper



First login to the Linux **workstation** in front of you.

Use your CARC username and password.

This is an “important step” so don’t let me move on until you have logged in

Logging into Hopper



```
ssh vanilla@hopper.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Logging into Hopper

Welcome to Hopper

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to <http://carc.unm.edu>, select the "New Users" menu and then click "Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the `chmod` command to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing, supported in part by the National Science Foundation, for providing the research computing resources used in this work."

Please send citations to publications@carc.unm.edu.

There are three types of slurm partitions on Hopper:

1) General - this partition is accessible by all CARC users.

2) Condo - preemtable scavenger queue available to all condo users. Your job must use checkpointing to use this queue or you will lose any work you have done if it is preempted by the partition's owner.

3) Named partitions - these partitions are available to condo users working under the grant/lab/center that purchased the associated hardware.

Type "qgrok" to get the status of the partitions.

Last login: Wed Jul 27 17:46:13 2022 from 129.24.246.68
mfricke@hopper:~ \$



Technology, IT etc.

SLURM

means

Simple Linux Utility for Resource
Management

by acronymsandslang.com

ENJOY

Slurmm
SODA

IT'S HIGHLY ADDICTIVE!

VOTED **#1** SOFT DRINK OF THE 31ST CENTURY!



SLURMS MCKENZIE

```
[vanilla@hopper~] $ qgrok
queues      free    busy    offline  jobs    nodes  CPUs  CPUs/node  Memory/node  time_limit  CPU_limit  RAM_limit
-----
general     3       7       0        5       10     320    32         93G         2-00:00:00  128        380000M
debug       1       1       0        1        2      64     32         93G         4:00:00     8          25000M
```

Compute Nodes are
divided into partitions
with different limits


```
[vanilla@hopper~] $ qgrok
queues      free    busy    offline  jobs    nodes  CPUs  CPUs/node  Memory/node  time_limit  CPU_limit  RAM_limit
-----
general      3       7       0        5       10     320    32         93G         2-00:00:00  128       380000M
debug        1       1       0        1        2      64     32         93G         4:00:00     8         25000M
```

The general partition
has high limits so you
generally have to wait
a little bit to run.

For debugged jobs.

```
[vanilla@hopper~] $ qgrok
queues      free    busy    offline  jobs    nodes  CPUs  CPUs/node  Memory/node  time_limit  CPU_limit  RAM_limit
-----
general     3       7       0        5       10     320    32         93G         2-00:00:00  128       380000M
debug       1       1       0        1        2      64     32         93G         4:00:00     8         25000M
```

**The debug partition
has low limits so
generally there is no
wait.**

**For debugging jobs.
For interactive jobs.**

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up        4:00:00    2    idle hopper[011-012]
```

sinfo reports information about
partitions

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       4:00:00    2    idle hopper[011-012]
```

The debug queues are intended
for testing your programs.

And for interactive jobs.


```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up        4:00:00    2    idle hopper[011-012]
```



Name

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       4:00:00    2    idle hopper[011-012]
```



You can run a “job” for up to 4 hrs.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       4:00:00    2    idle hopper[011-012]
```



There are two nodes in this partition.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       4:00:00    2   idle hopper[011-012]
```



The names of the nodes in the
partition


```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       4:00:00    2   idle hopper[011-012]
```



The names of the nodes in the
partition

```
[vanilla@hopper ~]$ sinfo --partition general
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
general*   up    2-00:00:00      9  alloc hopper[001-009]
general*   up    2-00:00:00      1   idle hopper010
```



Hopper001 through 009 are running jobs.

Hopper010 is waiting to be used.

```
[vanilla@hopper ~]$ hostname  
hopper  
[vanilla@hopper ~]$
```



Running on the Head Node.

The head node's name is "hopper".

```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname
```



```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man man  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ sinfo --all
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
general*	up	2-00:00:00	9	alloc	hopper[001-009]
general*	up	2-00:00:00	1	idle	hopper010
debug	up	4:00:00	2	idle	hopper[011-012]
condo	up	2-00:00:00	1	down*	hopper045
condo	up	2-00:00:00	3	mix	hopper[018-020]
condo	up	2-00:00:00	16	alloc	hopper[013-015,028-036,049-052]
condo	up	2-00:00:00	18	idle	hopper[016-017,021-027,037-044,053]
bugs	up	7-00:00:00	2	alloc	hopper[013-014]
pcnc	up	7-00:00:00	1	alloc	hopper015
pcnc	up	7-00:00:00	1	idle	hopper016
pathogen	up	7-00:00:00	1	idle	hopper017
tc	up	7-00:00:00	3	mix	hopper[018-020]
tc	up	7-00:00:00	2	alloc	hopper[029-030]
tc	up	7-00:00:00	5	idle	hopper[021-025]
gold	up	7-00:00:00	2	idle	hopper[026-027]
fishgen	up	7-00:00:00	1	alloc	hopper028
neuro-hsc	up	7-00:00:00	6	alloc	hopper[031-036]
neuro-hsc	up	7-00:00:00	8	idle	hopper[037-044]
cup-ecs	up	7-00:00:00	2	alloc	hopper[049-050]
tid	up	7-00:00:00	1	alloc	hopper051
biocomp	up	7-00:00:00	1	alloc	hopper052
chakra	up	7-00:00:00	1	idle	hopper053
pna	up	7-00:00:00	1	down*	hopper045

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Tell slurm to run a program
on a compute node...

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Run the program on a
compute node in the
debug partition.


```
[vanilla@hopper ~]$ srun --partition debug hostname
```



The program
to run.

```
[vanilla@hopper ~]$ srun --partition debug hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs,  
use the -G option in your submission script.  
hopper011
```

```
[vanilla@hopper ~]$ queue
```

```
[vanilla@hopper ~]$ queue
```

JOBID	PARTITION	NAME	USER	ST	TIM
4314	general	PRE	erowland	PD	0:00
4315	general	PRE	erowland	PD	0:00
4317	general	PRE	erowland	PD	0:00
4318	general	PRE	erowland	PD	0:00

**PD means programs
that are waiting their
turn.**

Shows you what the slurm scheduler is doing right now.

Here we can see that user 'erowland' has a lot of programs waiting to run.

[illegible]

```
[vanilla@hopper ~]$ queue
```

The reason these jobs are not running is that 'erowland' is already using the maximum number of CPUs they are allowed.

```
[vanilla@hopper ~]$ squeue -t R --all
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
4405	condo	2ndMA	mfricke	R	1-07:48:30	6	hopper[031-036]
5208	condo	NN	kgu	R	5:48:49	1	hopper015
5210	condo	NN	kgu	R	6:30:13	1	hopper014
5209	condo	NN	kgu	R	6:31:13	1	hopper013
5206	condo	NN	kgu	R	6:32:13	1	hopper051
5207	condo	NN	kgu	R	6:32:13	1	hopper052
5205	condo	NN	kgu	R	6:32:43	1	hopper028
4595	cup-ecs	golConfi	aalasand	R	2-06:51:59	1	hopper050
4594	cup-ecs	golConfi	aalasand	R	2-06:52:03	1	hopper049
5120	general	jupyterh	jacobm	R	11:45:47	1	hopper007
4313	general	PRE	erowland	R	1:17:29	2	hopper[003-004]
5111	general	1stMA	mfricke	R	11:15:28	2	hopper[005-006]
5025	general	c2n	jxzuo	R	1:50	1	hopper001
5024	general	c2n	jxzuo	R	31:28	1	hopper002
5203	general	NN	kgu	R	6:37:50	1	hopper009
5201	general	NN	kgu	R	6:38:14	1	hopper008
4390	tc	UCsTpCyd	lepluart	R	2-15:18:18	3	hopper[018-020]
5198	tc	NN	kgu	R	6:40:19	1	hopper030
5196	tc	NN	kgu	R	6:40:31	1	hopper029

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.  
hopper011  
hopper011
```

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.  
hopper011  
hopper011
```

You ran two **copies** of your program.

ntasks is the number of copies to run.


```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

```
hopper011
```

You ran eight **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project
```

hopper011

hopper011

hopper011

You have not been allocated GPUs. To request GPUs, use the -G option in your submission script.

hopper011

hopper011

hopper011

hopper011

hopper011

By default, each task (copy of your program) is allowed to use one CPU.

Many programs are able to use more than one CPU at a time.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G option in your submission  
script.
```

```
hopper011
```

```
hopper011
```

**Here we are telling SLURM to
run 2 copies of our program
and let each copy of our
program use 2 CPUs.**

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2 --ntasks-per-node 4 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using  
latest project
```

```
hopper012
```

```
You have not been allocated GPUs. To request GPUs, use the -G option in  
your submission script.
```

```
hopper012
```

```
hopper011
```

```
hopper011
```

```
hopper012
```

```
hopper012
```

```
hopper011
```

```
hopper011
```

**Here we are telling SLURM to run 4
copies of our program on 2 different
compute nodes.**

**This is useful when our programs need
a bigger share of the compute node.**

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2  
--ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
hopper011  
You have not been allocated GPUs. To request GPUs, use  
the -G option in your submission script.  
hopper011  
hopper012  
hopper012
```

And we can combine all three.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname
```

```
srun: Account not specified in script or  
~/default_slurm
```

```
hopper012
```

```
hopper012
```

```
You have not been granted access to  
the -G option in this partition.
```

```
hopper011
```

```
Hopper011
```

**And we can specify how much
memory we want.**

**--mem 4G means give me 4
gigabytes of memory per node.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname
```

```
srun: Account not specified in script or
```

```
~/default_slurm
```

```
hopper012
```

```
hopper012
```

```
You have not been
```

```
the -G option is
```

```
hopper011
```

```
Hopper011
```

Why does all this matter?

The purpose of SLURM is to provide you the hardware your programs need.

So you have to understand what those requirements are really well.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
hopper012
hopper012
You have not been assigned the -G option in
hopper011
Hopper011
```

- 1) Can my program use multiple CPUs?
- 2) How much memory does my program need?
- 3) Can my program use multiple compute nodes (MPI*, GNU Parallel*)?
- 4) Can my program use GPUs?


```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account using latest project  
hopper012  
hopper012  
You have not been  
the -G option in  
hopper011  
Hopper011
```

This command is getting pretty long.

We can use **salloc** to avoid asking for the same resources every time we use **srun**.

```
[vanilla@hopper ~]$ salloc --partition debug --nodes 2  
--ntasks-per-node 2  
salloc: Account not specified in script or  
~/.default_slurm_account, using latest project  
salloc: Granted job allocation 5251  
salloc: Waiting for resource configuration  
salloc: Nodes hopper[011-012] are ready for job  
[vanilla@hopper ~]$
```

This command is getting pretty long.

We can use **salloc** to avoid asking for the same resources every time we use **srun**.

```
[vanilla@hopper ~]$ srun hostname
```

```
hopper012
```

```
hopper012
```

```
hopper011
```

```
You have not been allocated GPUs. To request GPUs, use  
the -G option in your submission script.
```

```
hopper011
```

```
[vanilla@hopper ~]$ srun hostname
```

```
hopper012
```

```
hopper011
```

```
hopper012
```

```
hopper011
```

```
[vanilla@hopper ~]$
```

Now we can use **srun** over and over without having to ask for a new hardware allocation each time.

```
[vanilla@hopper ~]$ exit
```

```
exit
```

```
salloc: Relinquishing job allocation 5251
```

Always type **exit** when you are done with the hardware.

Running salloc inside an allocation gets very confusing.

Interactive vs Batch Mode

Interactive Mode

- Everything so far has been interactive. You request hardware, run your program, and get the output on your screen right away.

Batch Mode

- Most programs at an HPC center are run in “batch” mode.
- Batch mode means we write a shell script that the SLURM scheduler runs for us. The script requests hardware just like we did with `salloc` and then runs the commands in the script.
- Whatever would have been written to the screen is saved to a file instead.

```
[vanilla@hopper ~]$ git clone https://lobogit.unm.edu/CARC/workshops.git
Cloning into 'workshops'...
remote: Enumerating objects: 132, done.
remote: Counting objects: 100% (75/75), done.
remote: Compressing objects: 100% (43/43), done.
remote: Total 132 (delta 33), reused 74 (delta 32), pack-reused 57
Receiving objects: 100% (132/132), 57.58 KiB | 3.60 MiB/s, done.
Resolving deltas: 100% (51/51), done.
```

Rather than make you write shell scripts lets just download some we wrote for this workshop...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H20.gjf
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run tree to see how the workshops directories are organized...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│   ├── data
│   │   ├── H2O.gjf
│   │   └── step_sizes.txt
│   └── slurm
│       ├── calc_pi_array.sh
│       ├── calc_pi_mpi.sh
│       ├── calc_pi_parallel.sh
│       ├── calc_pi_serial.sh
│       ├── gaussian.sh
│       ├── hostname_mpi.sh
│       ├── vecadd_hopper.sh
│       ├── vecadd_xena.sh
│       ├── workshop_example2.sh
│       ├── workshop_example3.sh
│       └── workshop_example.sh
└── README.md
```

Run **tree** to see how the workshops directories are organized...

The workshop files are divided into “code”, “slurm”, and “data” directories.


```
[vanilla@hopper ~]$ cd workshops/intro_workshop
[vanilla@hopper intro_workshop]$ pwd
/users/vanilla/workshops/intro_workshop
[vanilla@hopper intro_workshop]$ cat slurm/workshop_example1.sh
#!/bin/bash
#SBATCH --partition debug
#SBATCH --ntasks 4
#SBATCH --time 00:05:00
#SBATCH --job-name ws_example
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

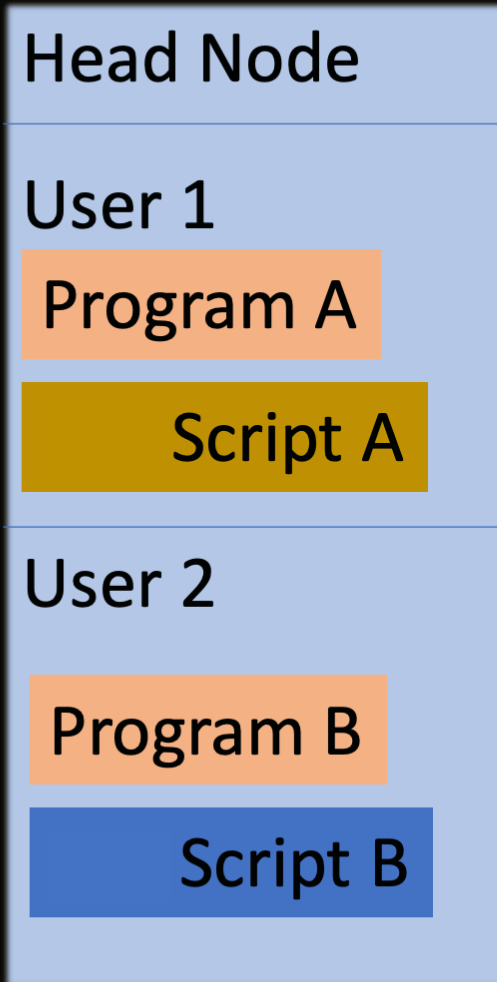
srun hostname
```

Let's take a look at the **workshop_example.sh** script in the slurm directory...

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh
sbatch: Account not specified in script or ~/.default_slurm_account,
using latest project
Submitted batch job 5252
[vanilla@hopper intro_workshop]$
```

We **submit** our slurm
shell script with the
sbatch command.

Workflow



Compute Node 01

Compute Node 02

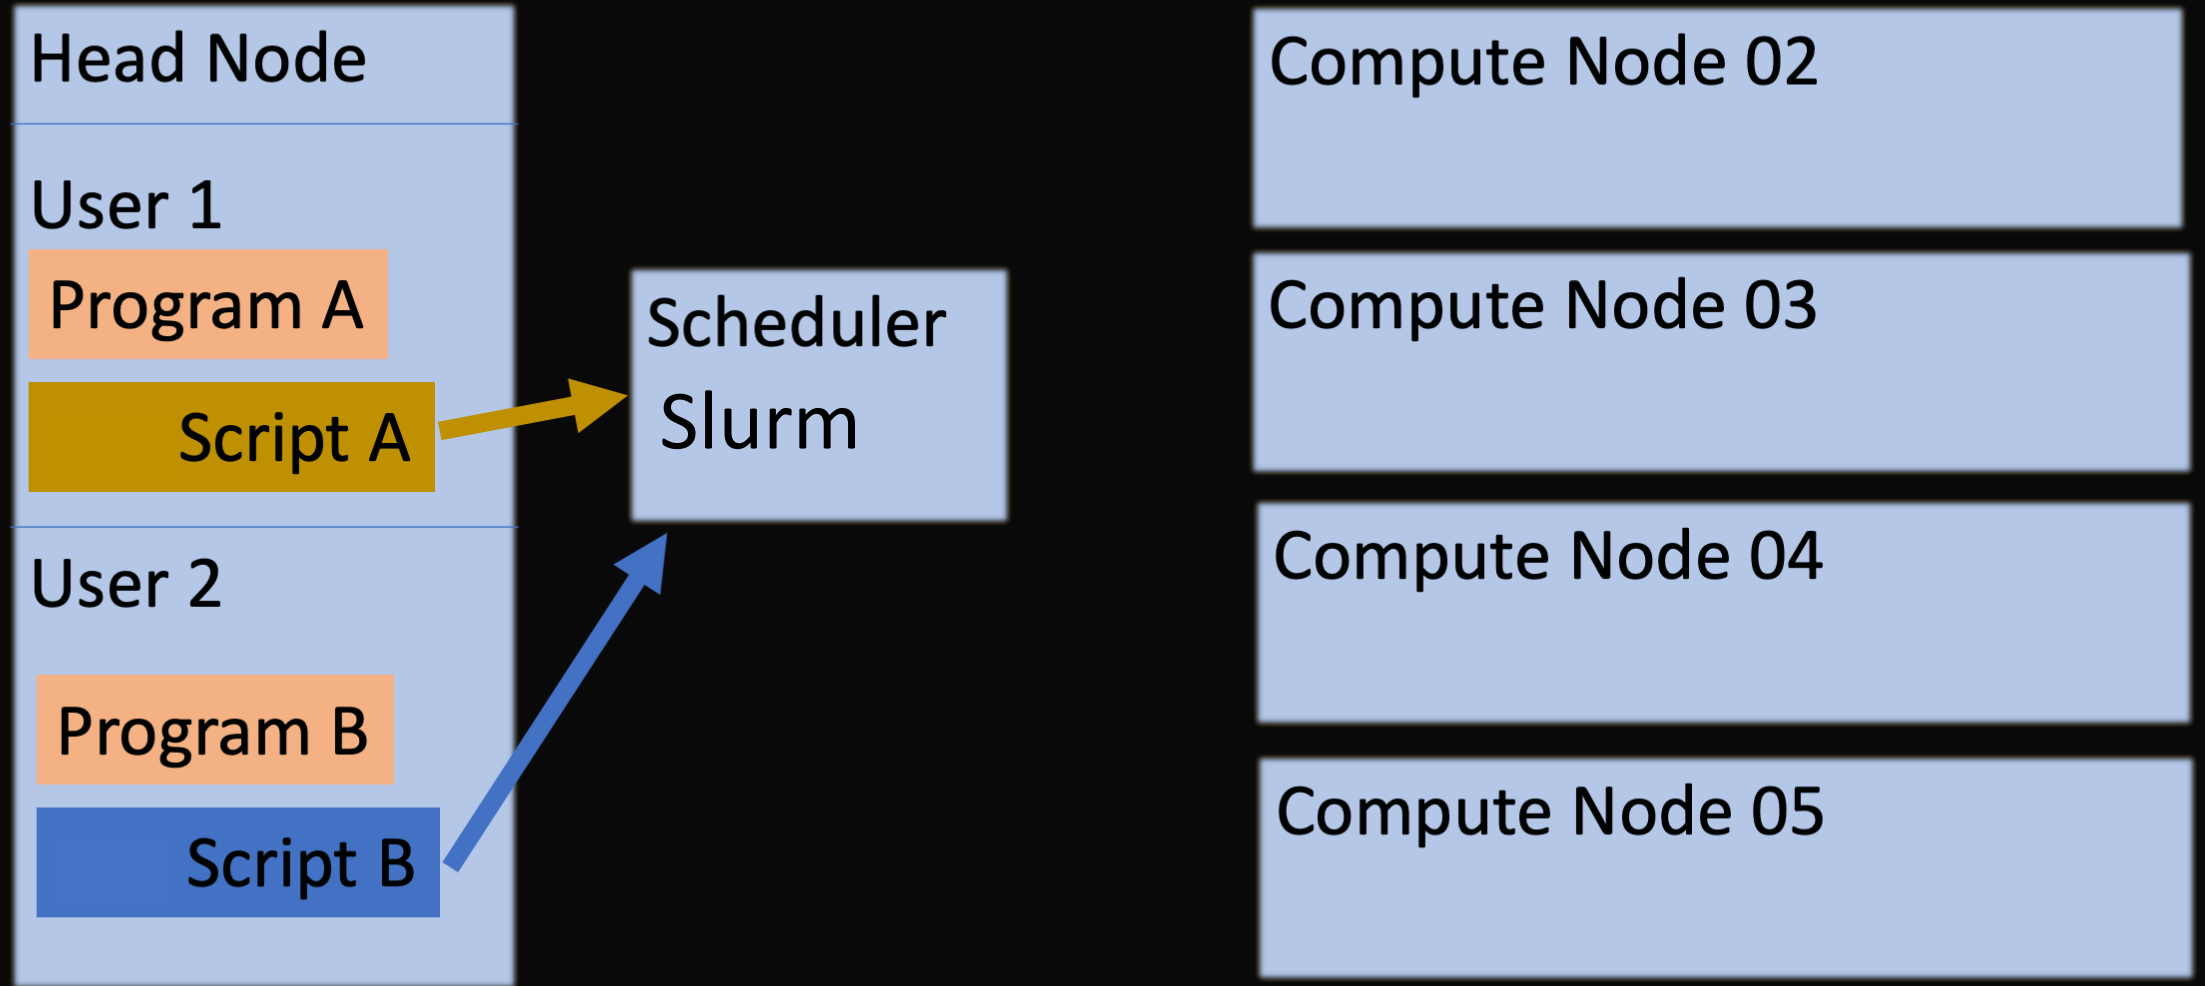
Compute Node 03

Compute Node 04

Compute Node 05

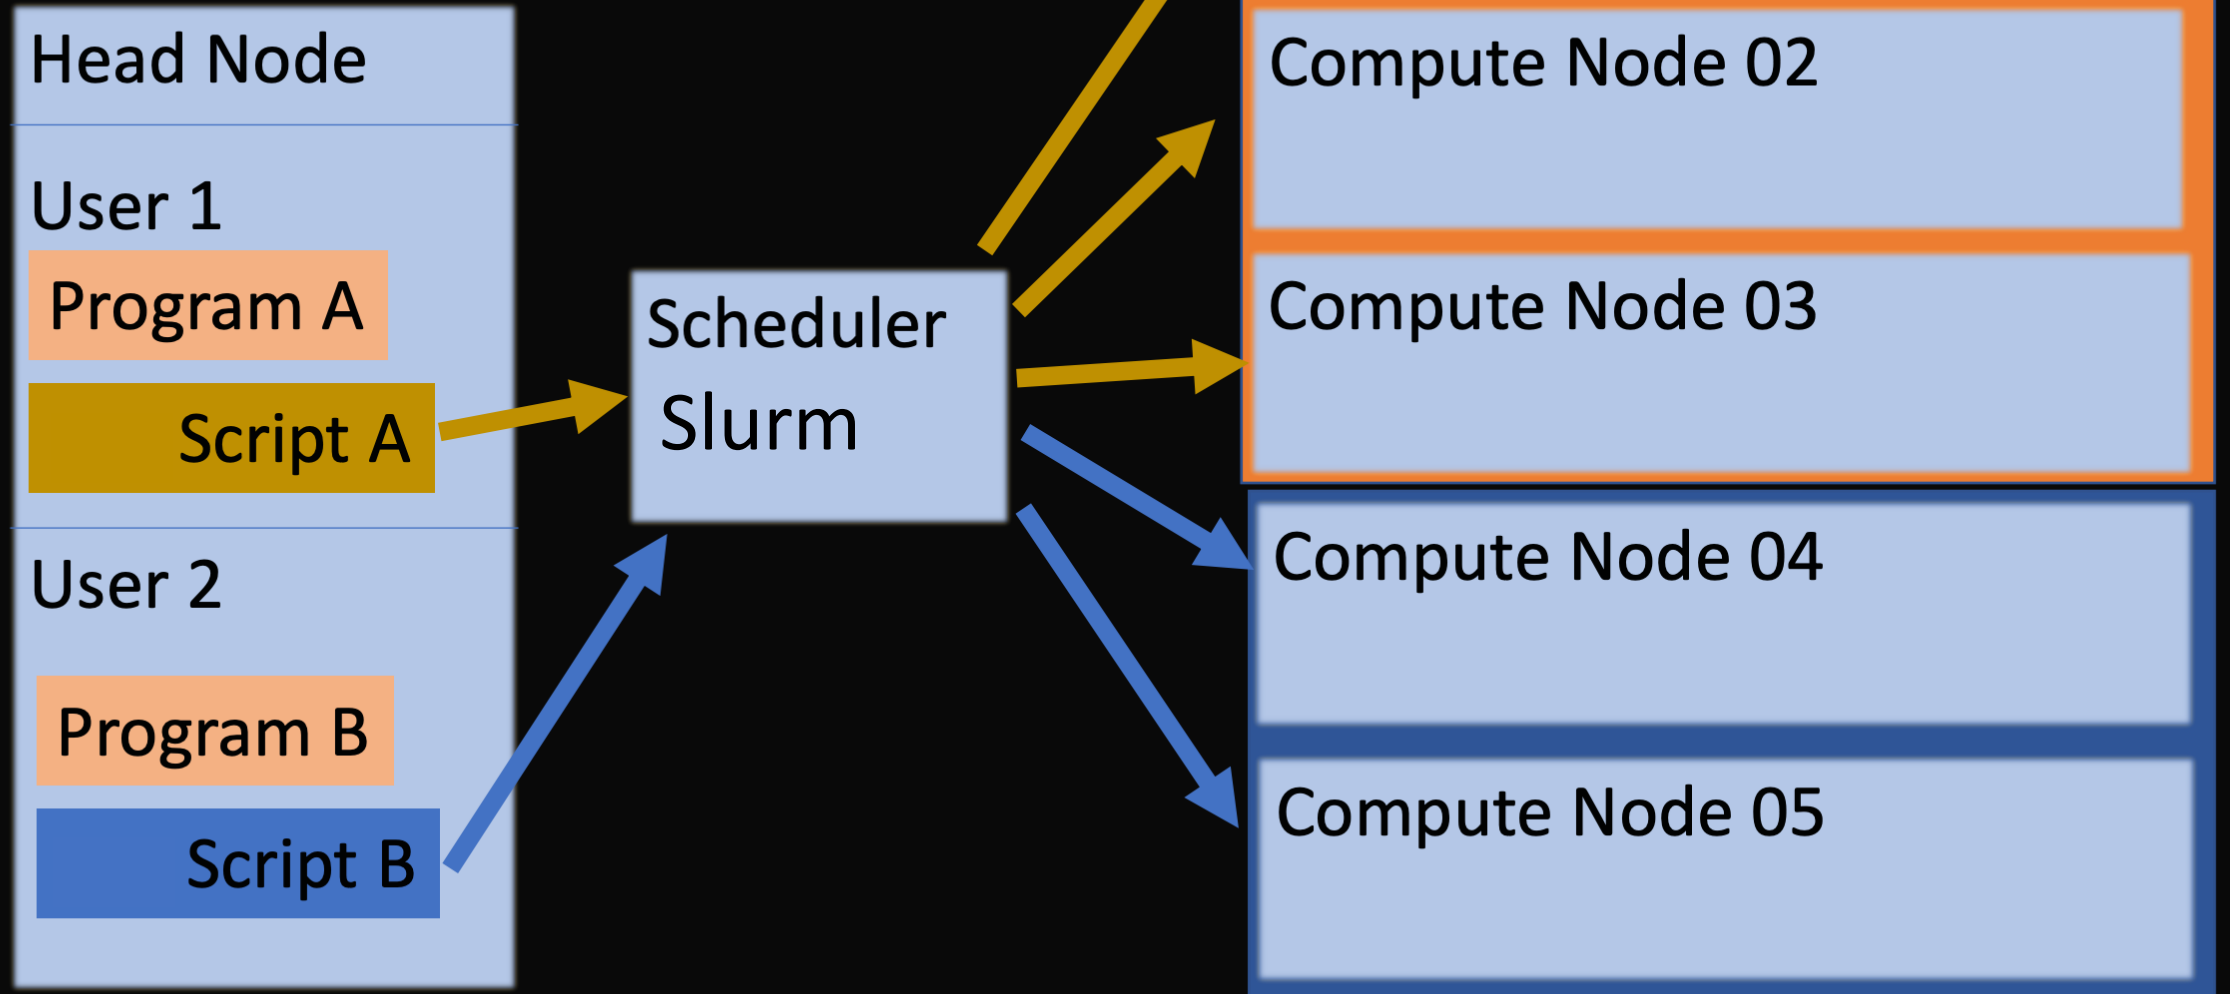
Shared filesystems – All nodes can access the same programs and write output

Workflow



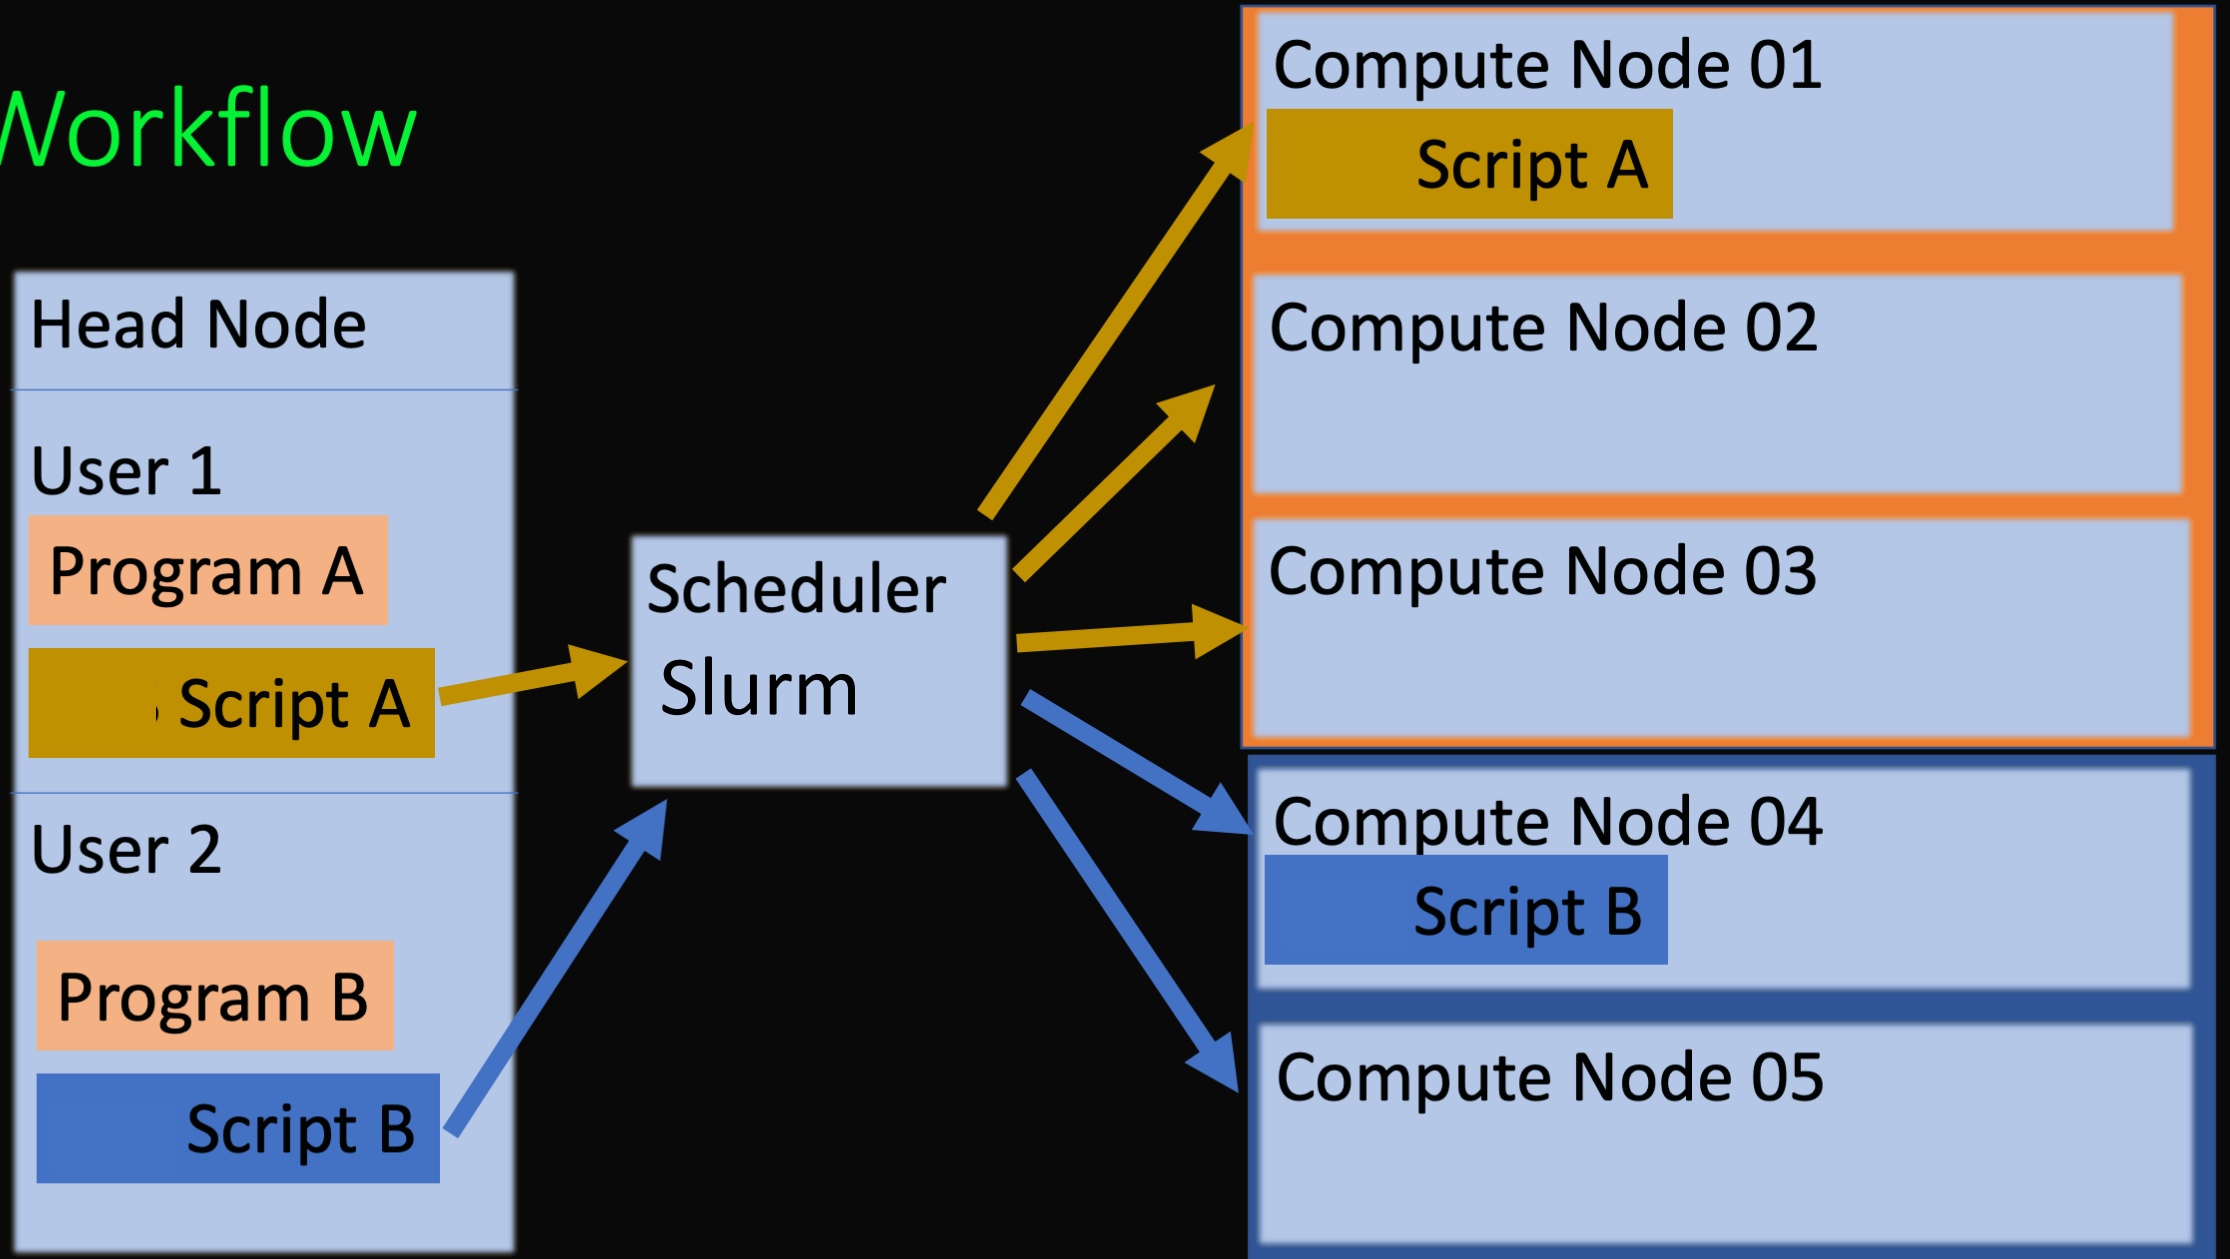
Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

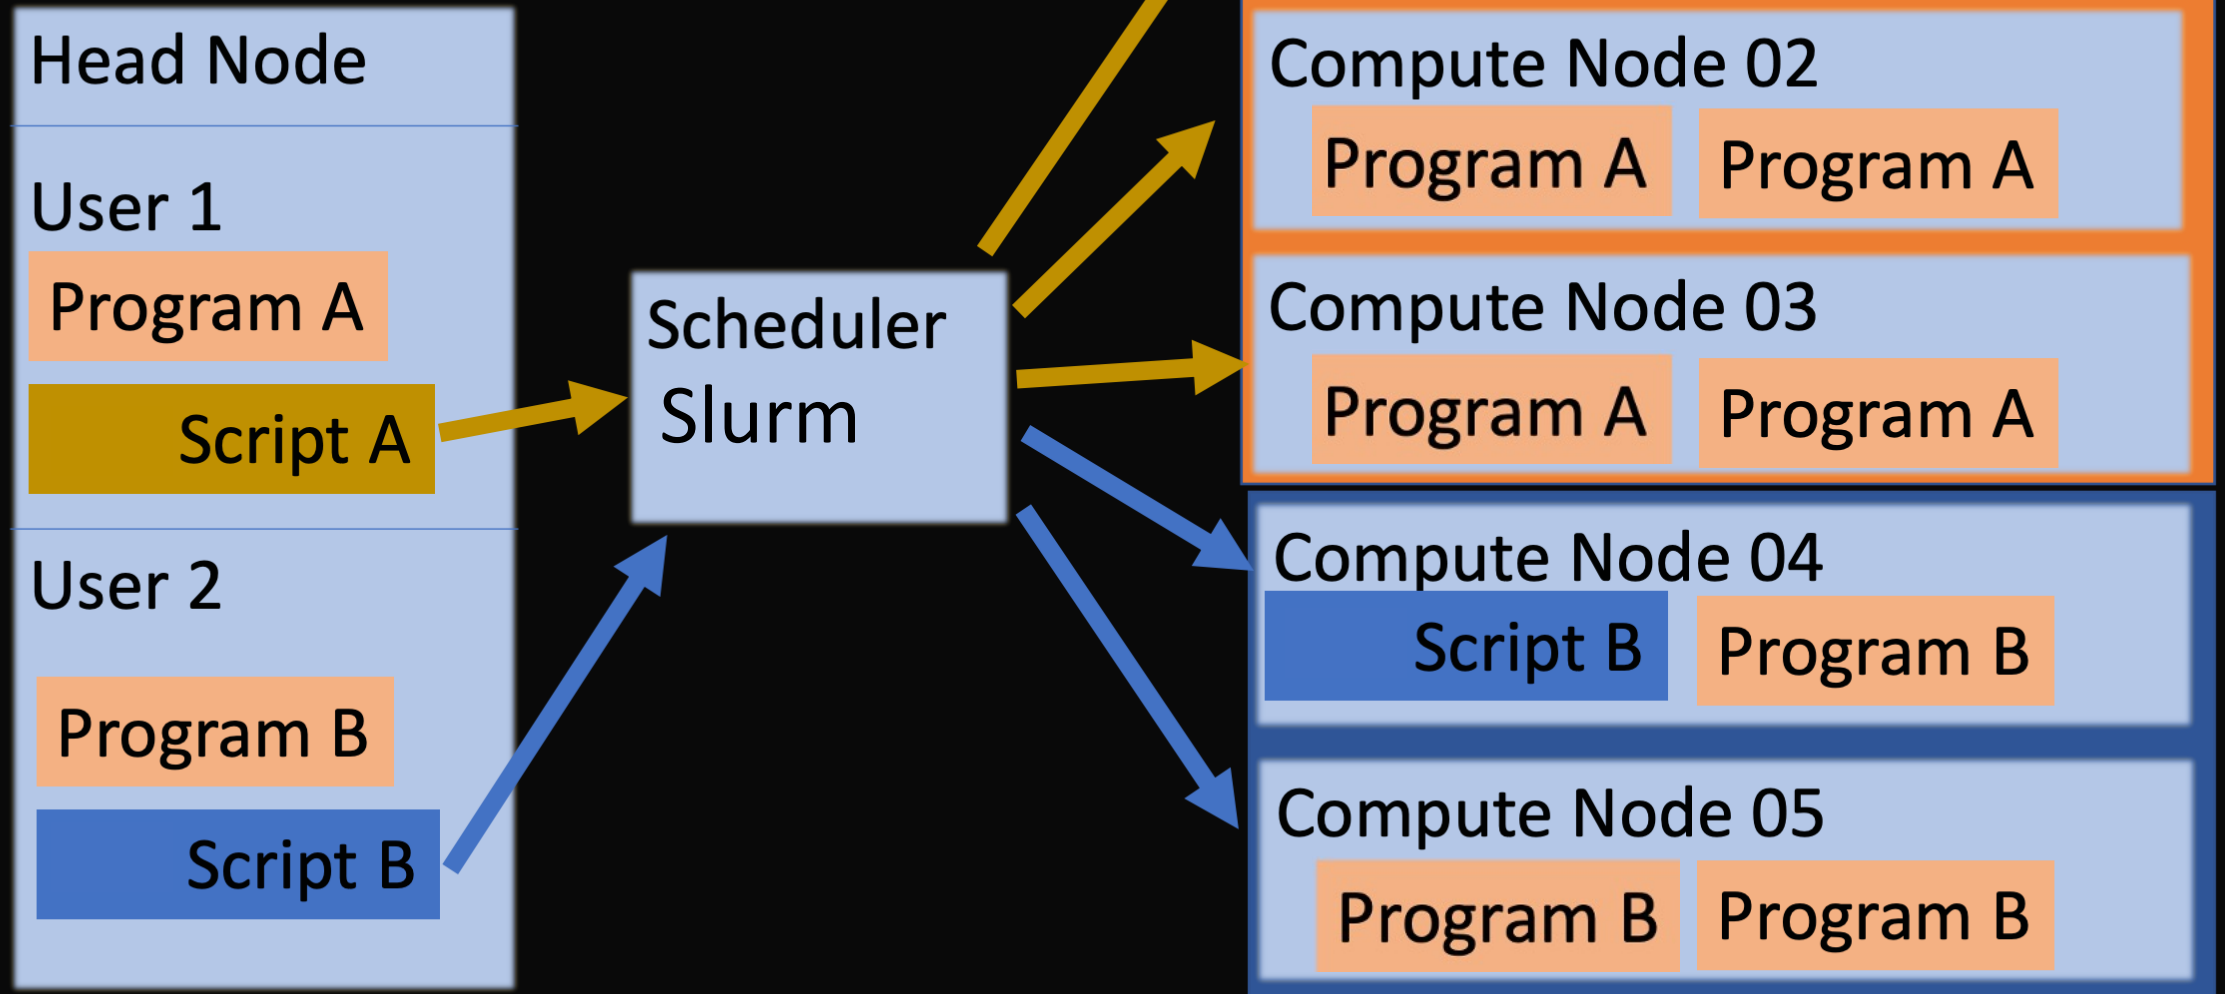
Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow

We need something in the script to run the program on all the nodes. E.g. `srun`.




Shared filesystems – All nodes can access the same programs and write output


```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```

The **hostname** command is very fast so everyone's job should finish in a few seconds.

When it is finished you will have a new file named `slurm-{your job id}.out`.





```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```



When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011  
hopper011  
hopper011  
hopper011
```

Serial Program to Calculate π



$$w = \frac{1}{n}$$
$$\pi \approx \sum_{i=0}^n \frac{4}{1 + \left(i + \frac{w}{2}\right)^2}$$

```
# A program that calculates pi using the area under a  
# Curve. The program checks the value of pi calculated  
# against the value provided by numpy
```

```
import time  
import sys  
import numpy as np # Value of PI to compare to
```

```
def Pi(num_steps): #Function to calculate pi  
    step = 1.0 / num_steps  
    sum = 0  
    for i in range(num_steps):  
        x = (i + 0.5) * step  
        sum = sum + 4.0 / (1.0 + x * x)  
        pi = step * sum  
    return pi
```

```
# Check that the caller gave us the number of steps to  
# use  
if len(sys.argv) != 2:  
    print("Usage: ", sys.argv[0], "<number of  
steps>")  
    sys.exit(1)
```

```
num_steps = int(sys.argv[1],10);
```

```
# Call function to calculate pi  
start = time.time() #Start timing  
pi = Pi(num_steps)  
end = time.time() # End timing
```

```
# Print our estimation of pi, the difference from  
numpy's value, and how long it took  
print("Pi = %.20f, (Diff=%.20f) (calculated in %f secs  
with %d steps)" %(pi, pi-np.pi, end - start,num_steps))  
sys.exit(0)
```

```
# A program that calculates pi using the area under a curve
# The program checks the value of pi calculated against the
# value provided by numpy
```

```
import time
```

```
import sys
```

```
import numpy as np # Value of PI to compare to
```

```
def Pi(num_steps): #Function to calculate pi
```

```
    step = 1.0 / num_steps
```

```
    sum = 0
```

```
    for i in range(num_steps):
```

```
        x = (i + 0.5) * step
```

```
        sum = sum + 4.0 / (1.0 + x * x)
```

```
    pi = step * sum
```

```
    return pi
```

```
# Check that the caller gave us the number of steps to use
if len(sys.argv) != 2:
    print("Usage: ", sys.argv[0], " <number of steps>")
    sys.exit(1)

num_steps = int(sys.argv[1],10);

# Call function to calculate pi
start = time.time() #Start timing
pi = Pi(num_steps)
end = time.time() # End timing

# Print our estimation of pi, the difference from numpy's value,
and how long it took
print("Pi = %.20f, (Diff=%.20f) (calculated in %f secs with %d
steps)" %(pi, pi-np.pi, end - start, num_steps))
sys.exit(0)
```

```
[vanilla@hopper intro_workshop]$ module load miniconda3  
[vanilla@hopper intro_workshop]$ conda create -n numpy numpy
```

Wait a while – introduce yourselves to your neighbor...

Conda allows you to install software into your home directory. In this case we need the numerical python libraries for calcPiSerial.py

Let's experiment with a program that does slightly more than print the hostname.

```
[vanilla@hopper intro_workshop]$source activate numpy  
[vanilla@hopper intro_workshop]$srun --partition debug python  
code/calcPiSerial.py 10  
srun: Using account 2016199 from ~/.default_slurm_account  
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.  
Pi = 3.14242598500109870940, (Diff=0.000833333141130559341)  
(calculated in 0.000005 secs with 10 steps)
```

**Activate the numpy environment and
Run calcPiSerial.py on a compute node.**

**For our example program the more steps it takes the
more accurate it is, but the longer it takes.**

```
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh  
sbatch: Using account 2016199 from ~/.default_slurm_account
```

```
Submitted batch job 5263
```

```
vanilla@hopper:~/workshops/intro_workshop$ squeue -me
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
5263	debug	calc_pi_	vanilla	R	0:44	1	hopper011

Edit `slurm/calc_pi_serial.sh`.


Change the email address to your address and submit the script.

Then enter `squeue --me` to see the job status.

Take a look at the job output.

Parallelism – Embarrassingly Parallel

- Embarrassingly parallel (Cleve Moler) are problems that are really really easy to speed up with more CPUs.
- The most common example is that you have a program that runs in serial and takes some input file, processes it, and produces some output.
- The problem is that you have 1,000 of the input files and want to run your program on each one.



Parallelism – Embarrassingly Parallel

This is “embarrassing”
because all you have to do
is run 1,000 copies of your
program on 1,000 CPUs
each with a different input
file and you are done.



```
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh
sbatch: Using account 2016199 from ~/.default_slurm_account
Submitted batch job 5263
vanilla@hopper:~/workshops/intro_workshop$ squeue -me
JOBID PARTITION      NAME      USER ST      TIME  NODES NODELIST(REASON)
5263      debug calc_pi_  vanilla R        0:44      1 hopper011
```

Submit the array script.

Then enter **squeue --me** to see the job status.


Take a look at the job output. (How many output files do you have?)

Parallelism – Coupled Parallelism

- Coupled problems are those where the CPUs need to work together to solve a problem by communicating with each other.
- Many commercial and research programs designed to run on HPC systems like CARC use a library called the message passing interface (MPI) to do this.
- We have written an MPI version of our python pi calculator to demonstrate.




Serial Program to Calculate π



$$w = \frac{1}{n}$$
$$\pi \approx \sum_{i=0}^n \frac{4}{1 + \left(i + \frac{w}{2}\right)^2}$$

Parallel Program to Calculate π



MPI: Message Passing Interface

When programs need to run on many processors but also communicate with one another.

Here the parallel version of calcPi needs to communicate the partial sums computed by each process so they can all be added up.

To communicate we will use the MPI library:

```
module load minconda3
```

```
conda create -n mpi_numpy mpi mpi4py numpy
```



```

import time
import sys
import numpy as np # Value of PI to compare to

##### SETUP MPI - START #####
from mpi4py import MPI      #Import the MPI library
comm = MPI.COMM_WORLD      #Communication framework
root = 0                   #Root process
rank = comm.Get_rank()     #Rank of this process
num_procs = comm.Get_size() #Total number of processes
##### END #####

#Distributed function to calculate pi
def Pi(num_steps):
    step = 1.0 / num_steps
    sum = 0
    for i in range(rank, num_steps, num_procs): # Divide
processes
        x = (i + 0.5) * step
        sum = sum + 4.0 / (1.0 + x * x)
    mypi = step * sum

    # Get that partial sums from all the processes, add them up, and give
to the root process
    pi = comm.reduce(mypi, MPI.SUM, root)
    return pi

```

```

#Main function
# Check that the caller gave us the number of steps to use
if len(sys.argv) != 2:
    print("Usage: ", sys.argv[0], " <number of steps>")
    sys.exit(1)

num_steps = int(sys.argv[1],10);

#Broadcast number of steps to use to the other processes
comm.bcast(num_steps, root)

# Call function to calculate pi
start = time.time() #Start timing
pi = Pi(num_steps) # Call the function that calculates pi
end = time.time() # End timing

# If we are the root process then print our estimation of pi,
# the difference from numpy's value, and how long it took
print("Pi = %.20f, (Diff=%.20f) (calculated in %f secs with %d
steps)" %(pi, pi-np.pi, end - start, num_steps))

```



```
##### SETUP MPI - START #####  
from mpi4py import MPI      #Import the MPI library  
comm = MPI.COMM_WORLD       #Communication framework  
root = 0                    #Root process  
rank = comm.Get_rank()      #Rank of this process  
num_procs = comm.Get_size() #Total number of processes  
##### END #####
```

#Distributed function to calculate pi

```
def Pi(num_steps):
```

```
    step = 1.0 / num_steps
```

```
    sum = 0
```

Divide sum among processes

```
    for i in range(rank, num_steps, num_procs):
```

```
        x = (i + 0.5) * step
```

```
        sum = sum + 4.0 / (1.0 + x * x)
```

```
    mypi = step * sum
```

Get that partial sums from all the processes, add them up,

and give to the root process

```
    pi = comm.reduce(mypi, MPI.SUM, root)
```

#Main function

<snip>

```
num_steps = int(sys.argv[1],10);
```

#Broadcast number of steps to use to the other processes

```
comm.bcast(num_steps, root)
```

Call function to calculate pi

```
start = time.time() #Start timing
```

```
pi = Pi(num_steps) # Call the function that calculates pi
```

```
end = time.time() # End timing
```

If we are the root process then print our estimation of pi,

the difference from numpy's value, and how long it took

```
print("Pi = %.20f, (Diff=%.20f) (calculated in %f secs with %d steps)"
```

```
%(pi, pi-np.pi, end - start, num_steps))
```

```
#!/bin/bash
#SBATCH --partition debug
#SBATCH --nodes 2
#SBATCH --ntasks-per-node 4
#SBATCH --time 00:05:00
#SBATCH --job-name calc_pi_mpi
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

module load miniconda3
source activate mpi_numpy

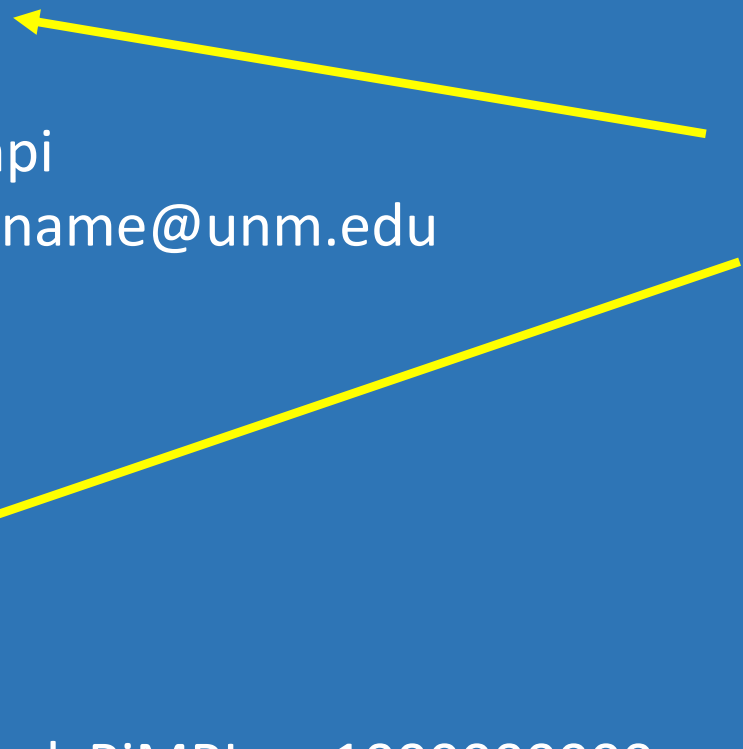
cd $SLURM_SUBMIT_DIR
srun --mpi=pmi2 python code/calcPiMPI.py 1000000000
```

sbatch slurm/calc_pi_mpi.sh

```
#!/bin/bash
#SBATCH --partition debug
#SBATCH --nodes 2
#SBATCH --ntasks-per-node 4
#SBATCH --time 00:05:00
#SBATCH --job-name calc_pi_mpi
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

module load miniconda3
source activate mpi_numpy

cd $SLURM_SUBMIT_DIR
srun --mpi=pmi2 python code/calcPiMPI.py 1000000000
```



srun inside a script
Automatically uses
Whatever you
allocated with
#SBATCH

`sbatch slurm/calc_pi_mpi.sh`

```
srun --mpi=pmi2 python code/calcPiMPI.py 1000000000
```

srun understands MPI programs!

If you ever used `mpirun` or `mpiexec` you had to provide a lot of parameters to describe how many compute nodes you had and what their names are, etc.

But `srun` is part of SLURM so it already knows all that.

The only thing you have to specify is the communication library to use. In our case “`pmi2`”.

Experiment

Run `calc_pi_mpi.sh`.

Vary the number of tasks it uses.

Use `squeue` to monitor the state of your job,

Look at your output files.

What is the relationship between the number of tasks and how fast it calculates pi?

```
[vanilla@hopper intro_workshop]$ srun --pty bash
srun: Account not specified in script or
~/.default_slurm_account, using latest project
To request GPUs, add --gpus-per-node X or --gpus X, where X is
the desired number of GPUs.
Job 2061894 running on hopper006
[vanilla@hopper034 ~]$ ls
[vanilla@hopper034 ~]$ exit
[vanilla@hopper~]$ module load matlab
[vanilla@hopper~]$ srun --pty matlab
```



You now have an interactive matlab session through a pseudoterminal!

The "--pty" Option

"--pty" creates a pseudoterminal that lets you interact with a program through the keyboard


```
[vanilla@hopper~]$ module load matlab
[vanilla@hopper~]$ srun --pty matlab
srun: job 879337 queued and waiting for resources
srun: job 879337 has been allocated resources
MATLAB is selecting SOFTWARE_OPENGL rendering.
```

```
      < M A T L A B (R) >
      Copyright 1984-2023 The MathWorks, Inc.
      R2023a Update 5 (9.14.0.2337262) 64-bit (glnxa64)
      July 24, 2023
```

To get started, type doc.
For product information, visit www.mathworks.com.

>>



You now have an interactive matlab session
through a pseudoterminal!

The “--pty” Option

“--pty” creates a pseudoterminal that lets you
interact with a program through the keyboard

Useful Slurm Commands

<code>squeue --me --long</code>	shows information about jobs you submitted
<code>squeue --me --start</code>	shows when slurm expects your job to start
<code>scancel jobid</code>	Cancels a job
<code>scancel --u \$USER</code>	Cancels all your jobs
<code>sacct</code>	shows your job history
<code>seff jobid</code>	shows how efficiently the hardware was used